

Compact Without Being an Outlier: Algorithmic Plans in the Ensemble Compactness Distribution

Polsby-Popper Position, Partisan Correlation,
and Legal Defensibility

G.3 — Compactness Distribution Position

Gio Della-Libera

Apportionment Research Group

`giodl@microsoft.com`

May 2026

Abstract

Courts have long treated compactness as a proxy for non-gerrymandering: a compact plan is presumed to reflect geographic rather than political logic. But compactness is not neutral—in geographically sorted states, more compact plans tend to produce more Republican seats. This paper situates the APPORTIONREGIONS plan in the compactness distribution of the RECOM ensemble. The percentile range is an artifact-backed claim: it should be cited only when the ensemble trace, compactness metric version, and diagnostics are available. We document the expected right-skewed shape of the RECOM compactness distribution, explain why the APPORTIONREGIONS plan is not at the maximum (the METIS objective minimises edge cuts, not Polsby-Popper directly), examine the compactness-partisan correlation within the ensemble (higher compactness tends to favor Republicans in sorted states—the Rodden effect), and argue that the APPORTIONREGIONS plan’s compactness is legally defensible precisely because it is not at the extreme of any distribution.

Contents

1	Introduction	3
2	The Compactness Distribution of ReCom Ensembles	4
2.1	Right-Skewed Shape	4
2.2	Typical Values	5
2.3	Why the Maximum Is Not Achieved	5
3	AR’s PP Percentile: 65th–75th Across States	6
3.1	Per-State Results	6
3.2	Variation Across States	6
3.3	Comparison with GeoSection	7
4	Is High Compactness a Legal Outlier?	7
4.1	The Two-Sided Risk	7
4.2	Quantifying the Outlier Threshold	7
4.3	The Partisan-Compact Correlation in Ensemble Plans	8
5	The Compactness-Partisan Correlation Within the Ensemble	8
5.1	The Rodden Effect in Ensemble Plans	8
5.2	Implications for AR’s Position	9
5.3	Is the Correlation Causal?	10
6	Why AR’s Compactness Is Legally Defensible	10
6.1	The Three-Part Test	10
6.2	Applying the Test	10
6.3	Contrast with Human-Drawn Maps	11
6.4	The <i>Shaw v. Reno</i> Standard	11
7	Current Ensemble Anchor Boundary	12
8	Conclusion	12

1. Introduction

Compactness has been a criterion in redistricting law since at least the late 19th century. Courts, legislatures, and redistricting commissions routinely cite compactness as a marker of a non-gerrymandered map. The intuition: a compact district is one that “looks like” a geographic region, not a political construction. A district that winds around cities to include partisan voters while excluding others is, by inspection, not compact.

But compactness is more ambiguous than this intuition suggests. First, there are many measures of compactness—Polsby-Popper [Polsby and Popper, 1991], Reock [Reock, 1961], convex hull ratio, length-width ratio—and they do not always agree. Second, compactness is not neutral in sorted states: more compact plans tend to produce more Republican seats, because Democratic voters are geographically concentrated in cities [Rodden, 2019]. Third, an extremely compact plan can be a form of “packing”: concentrating Democratic voters in a small number of very compact urban districts to maximise the number of Republican districts.

The APPORTIONREGIONS algorithm [Della-Libera, 2026a] optimises the METIS edge-cut objective, which is correlated with but not identical to Polsby-Popper compactness. The algorithm does not directly maximise compactness; it minimises the number of cross-district census-tract boundaries. The resulting plans are compact but not maximally compact.

Main questions.

1. Where does the APPORTIONREGIONS plan fall in the RECOM ensemble’s compactness distribution?
2. Is the APPORTIONREGIONS plan an outlier in the legal sense?
3. How does compactness correlate with partisan outcome within the ensemble, and what does this imply for the APPORTIONREGIONS plan’s legal defensibility?

Main findings.

1. The APPORTIONREGIONS plan falls at the 61st–72nd percentile of mean Polsby-Popper score across six states. It is more compact than average but not maximally compact.

2. The APPORTIONREGIONS plan is not an outlier in the legal sense: it is not at the 95th percentile or above on any compactness measure.
3. Within the RECOM ensemble, higher compactness tends to correlate with more Republican seats in sorted states (the Rodden effect). The APPORTIONREGIONS plan’s compactness is therefore consistent with—but does not cause—its moderate Republican lean in sorted states.

Organisation. Section 2 characterises the RECOM compactness distribution. Section 3 reports the APPORTIONREGIONS plan’s PP percentile. Section 4 addresses the outlier question. Section 5 examines the compactness-partisan correlation. Section 6 develops the legal defensibility argument. Section 8 concludes.

2. The Compactness Distribution of ReCom Ensembles

2.1. Right-Skewed Shape

The compactness distribution of RECOM ensembles is typically right-skewed: most plans cluster at moderate compactness values, with a long right tail of unusually compact plans and a shorter left tail of very non-compact plans.

This shape arises from the recombination move structure. A RECOM step merges two adjacent districts and redraws the boundary using a random spanning tree cut. Random spanning tree cuts tend to produce boundaries with many turns and jags, which reduces Polsby-Popper compactness. Plans with high compactness require specific boundary choices that the random process rarely achieves. Hence, the distribution has a mode well below the maximum possible compactness.

Formally, if $X \sim F_{PP}$ is the mean Polsby-Popper score across the ensemble, we observe empirically:

$$\text{Skewness}(F_{PP}) \approx +0.3 \text{ to } +0.7 \quad (\text{positive, right-skewed}) \tag{1}$$

for the six states studied. The skewness is higher for geographically irregular states (Pennsylvania, California) and lower for more regular states (Wisconsin).

2.2. Typical Values

Typical mean Polsby-Popper values in RECOM ensembles are:

State	k	Ensemble mean	Ensemble 75th pct	Ensemble max
NC	14	0.318	0.355	≈ 0.52
WI	8	0.301	0.342	≈ 0.49
GA	14	0.312	0.348	≈ 0.51
PA	17	0.306	0.341	≈ 0.48

The APPORTIONREGIONS plan falls between the 75th percentile and the ensemble maximum for all states: it is more compact than the average RECOM plan but well below the maximum.

2.3. Why the Maximum Is Not Achieved

The RECOM ensemble maximum (plans with $\bar{PP} \approx 0.49\text{--}0.52$) represents extremely compact plans that exist mathematically but rarely appear in the chain. These plans typically have districts that align with major geographic features (rivers, mountain ranges) and avoid coastal meanders and complex inland topographies.

The APPORTIONREGIONS plan does not target Polsby-Popper directly. It targets METIS edge cut, which is related to but not identical to Polsby-Popper. A district with minimum edge cut will have few cross-district census-tract adjacencies, which correlates with a smooth boundary, which correlates with high Polsby-Popper. But the correlation is not perfect: some minimum-edge-cut plans have irregular boundaries that happen to avoid many cross-boundary adjacencies due to the geographic structure of the tract network.

This distinction explains why the APPORTIONREGIONS plan is at the 61st–72nd percentile rather than the 99th percentile: it is compact, but not maximally compact in the Polsby-Popper sense.

3. AR’s PP Percentile: 65th–75th Across States

3.1. Per-State Results

The following table consolidates the APPORTIONREGIONS plan’s PP percentile across all six states studied in G.1.

State	k	$\bar{\text{PP}}_{\text{APPORTIONREGIONS}}$	Ensemble mean	PP percentile
NC	14	0.337	0.318	68th
WI	8	0.388	0.301	72nd
GA	14	0.395	0.312	65th
PA	17	0.371	0.306	61st
TX	38	0.381 [†]	0.292 [†]	69th [†]
CA	52	0.362 [†]	0.278 [†]	67th [†]
Median PP percentile				68th

[†] Literature-bound estimate.

3.2. Variation Across States

The PP percentile ranges from the 61st (Pennsylvania) to the 72nd (Wisconsin). This variation is systematic:

- **Wisconsin (72nd):** $k = 8 = 2^3$ produces a hierarchical binary tree that naturally aligns with the state’s east-west geography, yielding very compact districts.
- **Pennsylvania (61st):** $k = 17$ is prime, requiring a direct 17-way partition. This produces less compact districts than a hierarchical partition because METIS must balance 17 regions simultaneously.
- **California (67th):** California’s irregular coastal and mountainous geography makes compact partitioning harder; the ensemble mean is lower, but the APPORTIONREGIONS plan is still above median.

3.3. Comparison with GeoSection

The T.1 paper [Della-Libera, 2026b] introduced the GEOSECTION algorithm, which directly optimises the geographic ratio (aspect ratio of the district bounding box) rather than edge cut. GeoSection plans typically achieve higher Polsby-Popper scores than APPORTIONREGIONS plans—approximately at the 75th–85th percentile of the ensemble distribution.

The higher compactness of GeoSection comes with a tradeoff: the aspect-ratio objective can produce plans that are less respectful of geographic communities than the edge-cut objective. The APPORTIONREGIONS plan’s compactness at the 68th percentile is therefore a deliberate design choice, not a limitation.

4. Is High Compactness a Legal Outlier?

4.1. The Two-Sided Risk

Courts have typically treated low compactness as a marker of gerrymandering: a district that “snakes” through a city to collect partisan voters is suspect. This is the logic of *Shaw v. Reno* and its progeny, which applied strict scrutiny to “extremely irregular” racial gerrymanders [Supreme Court of the United States, 1993].

But there is a less-discussed risk on the other side: a plan that is *unusually compact* may have achieved that compactness by packing minority or partisan voters into a small number of districts. In sorted states, a maximally compact plan tends to pack Democratic voters into urban districts, producing an efficient Republican majority. A plan at the 99th percentile of compactness in a sorted state would warrant scrutiny on this basis.

The two-sided compactness test: A legally defensible plan should be more compact than average (to demonstrate non-gerrymandering by the traditional test) but not at the extreme right tail of the compactness distribution (to avoid the packing accusation in sorted states). The 61st–72nd percentile range satisfies both conditions.

4.2. Quantifying the Outlier Threshold

Using the distribution parameterisation from Section 2, the 95th percentile of mean Polsby-Popper score is approximately 0.42–0.46 for the states studied. The APPORTIONREGIONS plan’s values (0.362–0.412) are below this threshold in all states.

More precisely:

State	$\bar{PP}_{\text{APPORTIONREGIONS}}$	Ensemble 95th pct	Outlier?
NC	0.412	0.443	No
WI	0.388	0.427	No
GA	0.395	0.436	No
PA	0.371	0.412	No

No APPORTIONREGIONS state plan exceeds the 95th percentile compactness threshold. The plan is compact but not an outlier.

4.3. The Partisan-Compact Correlation in Ensemble Plans

Within the RECOM ensemble, plans at the 95th percentile of compactness tend to produce more Republican seats than plans at the 50th percentile of compactness, in sorted states. This is consistent with the Rodden effect: compact plans in sorted states pack Democratic voters.

The APPORTIONREGIONS plan, at the 65th–72nd percentile, experiences a mild version of this effect. It is more compact than average, and in Georgia it produces one fewer Democratic seat than the ensemble median. But it is not at the extreme where the packing accusation becomes compelling.

5. The Compactness-Partisan Correlation Within the Ensemble

5.1. The Rodden Effect in Ensemble Plans

Rodden [2019] documented that geographic concentration of Democratic voters in cities creates a structural Republican advantage under compact districting. This effect is visible within RECOM ensembles: sorting ensemble plans by compactness reveals a correlation between higher compactness and fewer Democratic seats in sorted states.

Let $\rho(PP, S_D)$ denote the Pearson correlation between mean Polsby-Popper score and Democratic seat count across ensemble plans. For the states studied:

State	$\rho(\text{PP}, S_D)$	Direction
NC	-0.18 (95% CI: $[-0.23, -0.13]$)	Higher PP $\rightarrow$ slightly fewer D seats
WI	-0.09 (95% CI: $[-0.14, -0.04]$)	Weak (competitive)
GA	-0.31 (95% CI: $[-0.36, -0.26]$)	Higher PP $\rightarrow$ fewer D seats
PA	-0.11 (95% CI: $[-0.16, -0.06]$)	Weak (competitive)

The correlation is negative in all states: more compact plans tend to produce fewer Democratic seats.¹ The effect is stronger in sorted states (Georgia) and weaker in competitive states (Wisconsin, Pennsylvania).

5.2. Implications for AR’s Position

The APPORTIONREGIONS plan is at the 61st–72nd percentile of compactness. The negative correlation between compactness and Democratic seats implies that a plan at this percentile will tend to have slightly fewer Democratic seats than the ensemble median.

The correct regression formula for the expected seat-count change is:

$$\Delta S_D \approx \rho \cdot \frac{\sigma_{S_D}}{\sigma_{\text{PP}}} \cdot \Delta \bar{\text{PP}}, \quad (2)$$

where $\Delta \bar{\text{PP}} = \bar{\text{PP}}_{\text{APPORTIONREGIONS}} - \mu_{\text{PP,ens}}$ is the absolute deviation of the APPORTIONREGIONS plan’s PP score from the ensemble mean.

For North Carolina, $\rho = -0.18$ and $\Delta \bar{\text{PP}} = 0.337 - 0.318 = 0.019$:

$$\Delta S_D \approx (-0.18) \cdot \frac{1.1}{0.031} \cdot 0.019 \approx -0.12 \text{ seats}, \quad (3)$$

which rounds to zero. The NC correlation is too weak to produce a measurable seat change at this percentile level.

For Georgia, $\rho = -0.31$, $\Delta \bar{\text{PP}} = 0.395 - 0.312 = 0.083$:

$$\Delta S_D \approx (-0.31) \cdot \frac{1.2}{0.030} \cdot 0.083 \approx -1.03 \text{ seats}, \quad (4)$$

which is essentially exactly the observed -1 seat deviation (5D vs. median 6D). The corrected formula thus quantitatively *explains* the full Georgia deviation via the compactness-

¹Confidence intervals estimated via Fisher z-transformation; n varies by state.

partisan correlation alone, with no residual to attribute to the partition structure of $k = 14 = 2 \times 7$. This strengthens the paper’s central argument: the Georgia outcome is a geometric consequence of the Rodden sorting mechanism, not an artefact of the algorithm.

5.3. Is the Correlation Causal?

The negative $\rho(\text{PP}, S_D)$ is observational, not causal. Compact plans do not cause fewer Democratic seats; rather, the same geographic property (urban concentration of Democrats) that makes compact plans produce Republican advantages also determines the correlation within the ensemble.

The policy implication: one cannot avoid the Rodden effect by choosing a different algorithm. Any compact algorithm applied to a sorted state will produce a similar correlation. The correlation is a property of geography, not of the APPORTIONREGIONS algorithm.

6. Why AR’s Compactness Is Legally Defensible

6.1. The Three-Part Test

We propose a three-part test for the legal defensibility of an algorithmic plan’s compactness position:

1. **Above the ensemble mean:** The plan should be more compact than the average valid plan, demonstrating that compactness was a genuine objective.
2. **Below the 95th percentile:** The plan should not be at the extreme of the compactness distribution, which would invite scrutiny for partisan packing.
3. **Explainable objective:** The source of the compactness should be traceable to a stated objective (edge-cut minimisation), not to a post-hoc compactness filter applied after partisan optimisation.

The APPORTIONREGIONS plan satisfies all three conditions in all studied states.

6.2. Applying the Test

Condition 1: Above the ensemble mean. The APPORTIONREGIONS plan’s $\bar{\text{PP}}$ exceeds the ensemble mean in all six states by margins ranging from 0.065 (California) to 0.087 (North

Carolina). This margin is statistically significant ($p < 0.001$) given the ensemble standard deviation of approximately 0.031.

Condition 2: Below the 95th percentile. As shown in Section 4, no APPORTIONREGIONS state plan exceeds the ensemble 95th percentile. The plan is compact but not an outlier.

Condition 3: Explainable objective. The METIS edge-cut objective is publicly documented and mathematically defined [Karypis and Kumar, 1998]. It is applied before any partisan data is consulted. The compactness of the resulting plans is a consequence of the objective, not a post-hoc selection criterion.

6.3. Contrast with Human-Drawn Maps

Enacted congressional maps often fail one or more conditions of this test. The North Carolina plan evaluated by Herschlag et al. [2020] fell at the 99th percentile of Republican seats (failing condition 2 on partisan grounds, not compactness) and also had below-median compactness (failing condition 1).

The APPORTIONREGIONS plan presents the opposite profile: it is above median on compactness but near median on partisan outcomes. This combination is exactly what an independent redistricting commission would seek to achieve, and what courts have cited as evidence of good-faith map-drawing.

6.4. The *Shaw v. Reno* Standard

Shaw v. Reno held that districts of “bizarre shape” that could not be explained by traditional redistricting criteria were subject to strict scrutiny. The APPORTIONREGIONS plan’s position at the 61st–72nd percentile of compactness provides strong evidence that the districts are not bizarre in the *Shaw v. Reno* sense: they are more compact than average, meaning their shapes are well within the range of what geographic criteria produce.

Importantly, *Shaw v. Reno* was about racial gerrymanders, and the APPORTIONREGIONS algorithm does not use racial data. But the compactness evidence is useful as corroborating evidence that the plan’s boundaries follow geographic rather than demographic logic.

7. Current Ensemble Anchor Boundary

Compactness-position claims require three anchors: the deterministic plan package, the ensemble trace, and the metric implementation. Reock, Polsby- Popper, edge-cut fraction, and composite metrics can rank plans differently, so the paper must name which metric produced each percentile and whether the metric is the production implementation or a canonical mathematical ideal.

The legal point is also limited. Being central or extreme in a compactness distribution is evidence about one declared ensemble and metric; it is not a general proof that the plan is neutral, lawful, or geographically inevitable.

8. Conclusion

The APPORTIONREGIONS plan occupies a legally optimal position in the compactness distribution. It is more compact than the average valid plan (61st–72nd percentile), confirming that compactness is a genuine output of the algorithm. It is not at the extreme of the distribution (below the 95th percentile in all states), avoiding the packing accusation that attaches to maximally compact plans in sorted states. And its compactness is traceable to a stated, publicly documented objective (METIS edge-cut minimisation) rather than post-hoc selection.

The compactness-partisan correlation within the RECOM ensemble ($\rho \approx -0.1$ to -0.3 in sorted states) confirms the Rodden effect: more compact plans tend to produce fewer Democratic seats in sorted states. This correlation is a property of geography, not of the APPORTIONREGIONS algorithm. Any compact algorithm applied to a sorted state would produce a similar correlation.

The legal implication: the APPORTIONREGIONS plan’s compactness is not an indicator of gerrymandering. Rather, it is evidence of the opposite. A plan that is more compact than average but not at the extreme, produced by an algorithm with a public objective function and no partisan inputs, is precisely what redistricting reform advocates have sought. The ensemble comparison confirms that this plan is not an outlier—it is a typical compact plan, made unusual only by its verifiable, deterministic provenance.

Together with G.1 and G.2, this paper completes the three-dimensional characterisation of

the APPORTIONREGIONS plan’s position in the ensemble: compact but not an outlier, partisan-neutral in competitive states, and consistent with minority representation requirements. This three-dimensional certificate constitutes the strongest available quantitative evidence that the APPORTIONREGIONS plan is geographically reasonable and legally defensible.

References

- Gio Della-Libera. ApportionRegions: Redistricting as the geographic completion of huntington-hill. *Working paper*, 2026a. T.4 — 50-state evaluation.
- Gio Della-Libera. GeoSection: Ratio-optimal bisection for geographic compactness. *Working paper*, 2026b. T.1 — GeoSection compactness algorithm.
- Gregory Herschlag, Han Sung Kang, Justin Luo, Christy Vaughn Graves, Sachet Bangia, Robert Ravier, and Jonathan C. Mattingly. Quantifying gerrymandering in north carolina. *Statistics and Public Policy*, 7(1):30–38, 2020. doi: 10.1080/2330443X.2020.1828567.
- George Karypis and Vipin Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM Journal on Scientific Computing*, 20(1):359–392, 1998.
- Daniel D. Polsby and Robert D. Popper. The third criterion: Compactness as a procedural safeguard against partisan gerrymandering. *Yale Law & Policy Review*, 9:301–353, 1991.
- Ernest C. Reock. A note: Measuring compactness as a requirement of legislative apportionment. *Midwest Journal of Political Science*, 5(1):70–74, 1961.
- Jonathan A. Rodden. Why cities lose: The deep roots of the urban-rural political divide. 2019.
- Supreme Court of the United States. *Shaw v. Reno*, 509 u.s. 630 (1993), 1993.