

VRASection: Geographic Alignment Score for Minority Opportunity District Bisection

Gio Della-Libera

Draft — May 12, 2026

Abstract

Congressional redistricting algorithms that attempt to satisfy the Voting Rights Act (§2) by adding minority population as an explicit METIS constraint (“MetisVra”) encounter a practical scale problem in the companion B.3 experiments: minority-fraction values (0.0–1.0) and census population counts (10,000–200,000) differ by orders of magnitude, producing infeasible or degenerate partitions in several tested configurations [Della-Libera, 2026a].

We introduce **VRASection**, which abandons explicit minority constraints and instead exploits the geographic structure of *Gingles Prong 1*: legally cognizable minority communities are, by definition, sufficiently large and geographically compact. This geographic concentration is already encoded in the census-tract adjacency graph; VRASection makes it measurable through a *geographic alignment score* $A(\text{split}) = |\text{MVAP}_{\text{frac}}(L) - \text{MVAP}_{\text{frac}}(R)|$ layered on top of GeoSection’s isoperimetric ratio scan [Della-Libera, 2026b]. In the current BISECT implementation, the score is computed from a demographic proxy: tract minority fraction multiplied by tract population, loaded from `data/{year}/demographics`. True VAP/CVAP inputs are supported in the audit data model, but the shipped ratio-selection path should be read as a minority-population spatial proxy rather than a legal finding.

The modified selection objective is:

$$\text{score}(\text{ratio}) = \text{EC}_{\text{norm}} - w_{\text{vra}} \cdot A \cdot \max(\text{EC}_{\text{norm}}, 1)$$

where $w_{\text{vra}} = 0.40$ by default. At this weight, compactness is weighted 50% more heavily than alignment. That coefficient is useful evidence for a geography-first design record, but it is not a mechanical safe harbor under *Shaw v. Reno* or *Miller v. Johnson*; racial predominance remains a fact-specific legal inquiry. The alignment score is best understood as a bounded use of a race-associated demographic attribute to detect spatial concentration, not as an automatic Gingles or constitutional certification.

Applied to Alabama (7 districts, 27% Black VAP), VRASection selects a 2 : 5 first-level bisection, separating the Black Belt region from the northern and coastal areas, compared to GeoSection’s 3 : 4 split (near-equal geographic division). The 2 : 5 split more effectively concentrates minority population on one side for downstream minority-opportunity-district formation. Full final-plan comparison packages across 6 historically significant VRA states remain pending; Alabama’s first-level 2 : 5 result provides the lead case study.

VRASection requires no user-specified minority-fraction targets. It passes the MetisVra failure modes by design: there is no second METIS constraint and no scale mismatch. We analyse the algorithm’s legal risk boundaries, sensitivity to w_{vra} , and relationship to the *Callais v. Landry* evidence workflow.

1 Introduction

The Voting Rights Act of 1965 (§2) prohibits electoral practices that result in minority voters having less opportunity than other members to participate in the political process and elect representatives of their choice. In the redistricting context, Section 2 creates an affirmative obligation: where a minority community satisfies the three-part *Gingles* test [Thornburg v. Gingles, 1986], mapmakers must provide the community with a realistic opportunity to elect candidates of its choice.

Algorithmic redistricting systems face a structural tension here. A map drawn by minimum-edge-cut bisection makes no reference to racial or political data: it sees only boundary lengths and census populations. This neutrality is legally valuable — it defeats racial-predominance challenges less directly than a race-targeting method under *Shaw v. Reno* [Shaw v. Reno, 1993] — but it means the algorithm has no mechanism to detect or correct VRA shortfalls.

Why explicit minority constraints fail. The natural engineering response is to add minority population as a second METIS constraint, alongside population balance. B.3 [Della-Libera, 2026a] tested this approach (“MetisVra”) on five majority-minority states with historically significant VRA litigation. The reported results were damaging: MetisVra failed completely in 3 of 5 states (Alabama, Louisiana, South Carolina), achieving zero majority-minority (MM) districts across all 140 tested parameter configurations.

The failure mode is **constraint scale mismatch**. METIS’s multi-constraint solver treats all constraint dimensions with the same numerical tolerance. Population balance operates on vertex weights of 10,000–200,000 people, requiring $\pm 0.5\%$ precision (± 50 – ± 1000 people). Minority VAP fraction operates on vertex weights of 0.0–1.0, where a $\pm 0.5\%$ tolerance is ± 0.005 — effectively zero. The scale difference is 60–800 $\times$. METIS’s local search cannot simultaneously satisfy both constraints: relaxing the minority constraint to a useful tolerance (e.g., $\pm 50\%$) provides so little guidance that METIS ignores it; tightening it creates a numerical conflict that produces infeasible partitions.

VRASection: geographic alignment as tie-breaker. We propose a different approach grounded in Gingles Prong 1 itself. Gingles requires that the minority community be “sufficiently large *and geographically compact*” to form a majority in a single-member district [Thornburg v. Gingles, 1986]. Geographic compactness is not an obstacle to algorithmic redistricting: it is a property that is *already encoded in the census-tract adjacency graph*. A minority community that satisfies Prong 1 is, by definition, one whose tracts cluster together in the adjacency graph.

VRASection makes this clustering measurable without adding a METIS constraint. At the first bisection level, it evaluates each candidate ratio $i : (k - i)$ by a two-component score:

1. **Compactness:** the isoperimetrically normalised edge-cut $EC_{\text{norm}}(i)$ from GeoSection [Della-Libera, 2026b] (lower is better).
2. **Geographic alignment:** $A(\text{split}) = |MVAP_{\text{frac}}(L) - MVAP_{\text{frac}}(R)|$, where $MVAP_{\text{frac}}$ is the share of the current subgraph’s minority-population signal placed on each side of the bisection.

The combined selection score is:

$$\text{score}(\text{ratio}) = EC_{\text{norm}} - w_{\text{vra}} \cdot A \cdot \max(EC_{\text{norm}}, 1)$$

Lower scores are preferred. The alignment term *subtracts* from the compactness score: a ratio that places compact minority communities on one side of the cut receives a lower (better) score than an otherwise identical ratio with dispersed minority population.

Alabama: the lead case study. Applied to Alabama (7 districts, 27.5% Black VAP, $w_{\text{vra}} = 0.40$), VRASection selects a 2 : 5 first-level bisection, separating the Black Belt region (south-central Alabama, historically concentrated Black population) from the northern industrial corridor and Gulf Coast. GeoSection selects a 3 : 4 split (near-equal geographic division; see T.1 Table 1 for the confirmed AL GeoSection natural ratio). The 2 : 5 split better preserves the geographic concentration relevant to Gingles Prong 1 because it concentrates a larger fraction of Alabama’s Black population on one side, providing a geographic foundation for downstream MM district formation without imposing a district-level racial target.

The legal distinction. VRASection uses a race-associated demographic input, which raises legal risk that must be documented rather than waved away. The critical difference is in the function computed. MetisVra’s second constraint forces METIS to concentrate minority population above a user-specified threshold — race is a co-equal optimization target. VRASection’s alignment score measures the *spatial distribution* of any population that the user designates as “minority VAP”. The algorithm selects the bisection ratio that best exploits pre-existing geographic concentration in the user-provided demographic signal. At $w_{\text{vra}} = 0.40$, compactness is weighted $1/(1 - 0.40) = 1.67\times$ more heavily than alignment. That supports a geography-first record, but *Shaw*’s racial-predominance test remains fact-specific.

Contributions. This paper makes four contributions:

1. **VRASection algorithm** (§3): formal definition of the alignment score, modified objective function, and algorithm pseudocode, including legal analysis of how the design limits racial-predominance risk.
2. **Alabama case study** (§4): the 2 : 5 vs. 3 : 4 comparison as the lead empirical result, with analysis of the Black Belt geographic concentration that drives the split.
3. **Comparison framework** (§4): a six-state comparison table covering the historically significant VRA states, with cells marked for runs pending full pipeline comparison.
4. **Legal and sensitivity analysis** (§5): the Shaw predominance line as a function of w_{vra} , relationship to MetisVra failures, and comparison to the Callais evidence standard.

Paper organisation. Section 2 reviews related work on VRA compliance, gerrymandering jurisprudence, and geographic concentration. Section 3 defines VRASection formally. Section 4 presents the Alabama empirical result and the multi-state comparison framework. Section 5 discusses legal implications, w_{vra} sensitivity, and limitations. Section 6 concludes.

2 Related Work

Voting Rights Act Section 2 and the Gingles test. Section 2 of the Voting Rights Act of 1965, as amended in 1982, prohibits electoral practices that result in minority voters having less opportunity than other members to participate in the political process [United States Congress, 1965]. In *Thornburg v. Gingles* (1986), the Supreme Court established a three-part test for Section 2 vote dilution claims in redistricting [Thornburg v. Gingles, 1986]:

1. **Prong 1 (geographic compactness):** The minority group is sufficiently large and geographically compact to constitute a majority in a single-member district.
2. **Prong 2 (political cohesion):** The minority group is politically cohesive.
3. **Prong 3 (bloc voting):** The white majority votes sufficiently as a bloc to usually defeat the minority’s preferred candidate.

Gingles Prong 1 is the geographic predicate: it asks whether the minority community is compact enough to form a district. VRASection operationalises one proxy for Prong 1: the alignment score A measures geographic concentration in the supplied demographic signal. A threshold such as $A \geq 0.15$ can flag communities for review, but it is not itself a legal Prong 1 finding.

Shaw v. Reno and racial predominance. *Shaw v. Reno* (1993) and *Miller v. Johnson* (1995) established that racial considerations cannot be the “predominant factor” in district design [Shaw v. Reno, 1993, Miller v. Johnson, 1995]. A district is presumptively unconstitutional if it is so “bizarrely shaped” that it is unexplainable on grounds other than race. The predominance test asks whether the legislature subordinated traditional redistricting principles — compactness, contiguity, preservation of political subdivisions — to racial considerations.

VRASection’s w_{vra} parameter makes part of the design record quantitative: at $w_{\text{vra}} < 0.5$, the compactness component of the selection score exceeds the alignment component for any ratio. This supports a geography-first explanation, but it does not ensure the legal conclusion (Section 3.4). At $w_{\text{vra}} = 0.40$ (the proposed default), compactness is explicitly weighted 60 % against alignment’s 40 %.

Allen v. Milligan and the modern Section 2 standard. *Allen v. Milligan* (2023) reaffirmed the Gingles framework and held that Alabama’s 2021 congressional map, which produced only one Black-majority district despite Alabama’s 27 % Black VAP, violated Section 2 [Allen v. Milligan, 2023]. The Court’s 5–4 opinion confirmed that race-neutral algorithms do not insulate a map from Section 2 challenge if the resulting map dilutes minority voting

strength. Alabama is the lead case study for VRASection precisely because its geography — concentrated Black Belt population, dispersed northern industrial and coastal areas — provides a strong alignment signal that the algorithm should detect and act on.

Callais v. Landry and the Louisiana remedy. *Callais v. Landry* (2025) addressed Louisiana’s remedial maps following Section 2 litigation [Callais v. Landry, 2025]. The case illustrates the tension between Section 2 remedies and traditional redistricting criteria. VRASection is designed to produce plan candidates for this kind of review: the alignment score concentrates minority population without abandoning compactness or contiguity. Validation against Louisiana’s Section 2 remedy (2 Black-majority districts from $k = 6$) is a key experiment in the full empirical suite.

Geographic concentration and the Rodden thesis. Rodden [2019] documents the deep geographic sorting of American voters: Democrats are concentrated in dense urban cores; Republicans are distributed across suburban and rural areas. This sorting interacts with redistricting in two ways. For partisan purposes, minimum-edge-cut bisection tends to isolate urban cores, producing the caterpillar pathology identified in T.1 [Della-Libera, 2026b]. For VRA purposes, the Rodden thesis implies that minority communities (disproportionately Democratic) are already geographically concentrated in urban cores and specific rural corridors (e.g., the Black Belt, the Rio Grande Valley, the Mississippi Delta). VRASection exploits this concentration directly: states where minority communities have high geographic clustering will exhibit high alignment scores, and the algorithm’s tie-breaking will concentrate them on one side of the first bisection.

Chen and Rodden [2013] confirm that geographic sorting produces predictable partisan outcomes under compactness criteria. VRASection’s alignment score provides a parallel signal for minority communities: the geographic sorting that makes a community legally cognizable under Gingles Prong 1 is the same sorting that produces a high alignment score.

Compactness in the VRA context. Compactness serves a dual function in VRA jurisprudence. Under Gingles Prong 1, compactness is a predicate for coverage: only compact minority communities trigger Section 2’s affirmative obligation. Under Shaw, compactness is a limiting principle: racial predominance is established when compactness is sacrificed in service of racial targets. VRASection navigates this tension by making compactness the primary criterion (the EC_{norm} term in the objective) and alignment a bounded secondary term ($w_{\text{vra}} \leq 0.5$). The algorithm rewards natural geographic concentration relevant to Gingles Prong 1 while preserving compactness as the larger objective term relevant to Shaw.

MetisVra (B.3): the explicit constraint approach. B.3 [Della-Libera, 2026a] introduced MetisVra, which adds minority VAP as an explicit second METIS constraint ($n_{\text{con}}=2$). The approach fails in 3 of 5 tested states due to constraint scale mismatch: minority VAP fractions (0.0–1.0) and census population counts (10,000–200,000) differ by 60–800 $\times$, and METIS’s numerical solver cannot simultaneously satisfy both constraints with meaningful precision. Alabama (0% success across 140 configurations) and Louisiana (0%) are the clearest failures; South Carolina also fails completely. MetisVra succeeds only in states with

large k (Georgia, $k = 14$) or very high minority concentration (Mississippi, $k = 4$, 38.8% Black VAP).

VRASection addresses the MetisVra failure at the architectural level: there is no second METIS constraint, so there is no scale mismatch. The alignment score is computed post-hoc from the METIS output, not used as a constraint input.

GeoSection (T.1): the isoperimetric ratio scan. VRASection extends GeoSection [Della-Libera, 2026b], which introduced isoperimetrically normalised ratio scanning for recursive bisection. GeoSection’s key contribution — the $EC/\sqrt{\min(i, k-i)}$ normalisation that prevents the caterpillar pathology — carries over to VRASection unchanged. VRASection adds a single modification: the selection score at the first bisection level incorporates the alignment bonus. All subsequent recursion levels use standard GeoSection (no alignment), so that race is observed at most once in the bisection tree.

3 The VRASection Algorithm

3.1 Setup and Notation

Let $G = (V, E, w)$ be the state’s census-tract adjacency graph. Vertices V are census tracts; edges E connect tracts sharing a common boundary; edge weights $w(e) > 0$ are shared boundary lengths in metres. Vertex weights $\mathbf{p} \in \mathbb{Z}_{>0}^{|V|}$ are census tract populations. Let k be the number of congressional districts allocated to this state. Let $\mathbf{m} \in \mathbb{R}_{\geq 0}^{|V|}$ be a per-tract minority population signal. In the current BISECT ratio-selection implementation, m_v is the tract’s minority fraction multiplied by tract population; this is a spatial proxy, not necessarily true VAP or CVAP. When true VAP/CVAP is available, the same formulas can be applied to those counts.

Definition 1 (Feasible bisection). *A feasible bisection of G with ratio $i : (k-i)$, $1 \leq i \leq \lfloor k/2 \rfloor$, is a partition (L, R) with $L \cup R = V$, $L \cap R = \emptyset$, and $\sum_{v \in L} p_v / \sum_{v \in V} p_v \in [i/k \cdot (1 - \delta), i/k \cdot (1 + \delta)]$ for balance tolerance δ (congressional default: $\delta = 0.005$). The edge-cut of the bisection is $EC(L) = \sum_{e \in \partial(L)} w(e)$ where $\partial(L)$ is the set of edges crossing the L/R boundary.*

Definition 2 (Isoperimetrically normalised edge-cut). *For ratio $i : (k-i)$, the normalised edge-cut is:*

$$EC_{norm}(i) = \frac{\min_{(L,R) \text{ feasible at } i:(k-i)} EC(L)}{\sqrt{\min(i, k-i)}}$$

The $\sqrt{\min(i, k-i)}$ denominator is GeoSection’s isoperimetric correction [Della-Libera, 2026b], which ensures that all ratios compete on equal footing.

3.2 Geographic Alignment Score

Definition 3 (Minority VAP fraction). *For a subset $S \subseteq V$, the minority VAP fraction is:*

$$MVAP_{frac}(S) = \frac{\sum_{v \in S} m_v}{\sum_{v \in V} m_v}$$

where the denominator is over the full vertex set V of the current subgraph (not just S). This measures what fraction of the subgraph’s total minority population lands on side S of the bisection.

Definition 4 (Alignment score). *For a feasible bisection (L, R) of G , the alignment score is:*

$$A(L, R) = |\text{MVAP}_{\text{frac}}(L) - 0.5| \times 2 = \left| \frac{\sum_{v \in L} m_v}{\sum_{v \in V} m_v} - 0.5 \right| \times 2$$

Equivalently, $A(L, R) = |\text{MVAP}_{\text{frac}}(L) - \text{MVAP}_{\text{frac}}(R)|$.

The alignment score has a natural interpretation. $A = 0$ means minority population is split exactly 50/50 between L and R — the bisection is *dispersive*: it divides the minority community across both sides. $A = 1$ means all minority population is on one side — the bisection is *concentrating*: it places the entire minority community on one side. A high alignment score indicates that the state has a natural geographic boundary separating high-minority from low-minority areas, and that the bisection at ratio $i : (k - i)$ falls along this boundary.

3.3 Modified Selection Objective

VRASection modifies GeoSection’s ratio selection criterion by combining the normalised edge-cut with the alignment score.

Definition 5 (VRASection selection score). *For ratio $i : (k - i)$ with best-seed bisection (L^*, R^*) at that ratio, the VRASection selection score is:*

$$\text{score}(i) = \text{EC}_{\text{norm}}(i) - w_{\text{vra}} \cdot A(L^*, R^*) \cdot \max(\text{EC}_{\text{norm}}(i), 1)$$

where $w_{\text{vra}} \in [0, 1]$ is the alignment weight (default 0.40). The ratio i^* selected by VRASection minimises this score:

$$i^* = \arg \min_{1 \leq i \leq \lfloor k/2 \rfloor} \text{score}(i)$$

The $\max(\text{EC}_{\text{norm}}(i), 1)$ factor ensures the alignment bonus scales with the compactness cost: a ratio with high EC gets a proportionally larger bonus for good alignment, so alignment can make a moderately expensive ratio competitive with a very cheap but dispersive ratio. For typical states where $\text{EC}_{\text{norm}} \gg 1$ (edge cuts are in the hundreds of kilometres), this reduces to:

$$\text{score}(i) \approx \text{EC}_{\text{norm}}(i) \cdot (1 - w_{\text{vra}} \cdot A(L^*, R^*))$$

which has a clean interpretation: the effective compactness cost is discounted by the alignment bonus.

Proposition 1 (VRASection reduces to GeoSection at $w_{\text{vra}} = 0$). *When $w_{\text{vra}} = 0$, $\text{score}(i) = \text{EC}_{\text{norm}}(i)$ for all i , and VRASection selects the same ratio as GeoSection.*

Proof. Immediate from Definition 5: the alignment term vanishes when $w_{\text{vra}} = 0$. □

Proposition 2 (Pure geographic concentration at $w_{vra} = 1$). *When $w_{vra} = 1$, the selection score is $\text{score}(i) = \text{EC}_{\text{norm}}(i) \cdot (1 - A(L^*, R^*))$, and the algorithm maximises geographic alignment subject to compactness being non-negative. Among ratios where $A = 1$ (perfect concentration), the ratio with smallest EC_{norm} is selected.*

Proof. At $w_{vra} = 1$ and for $\text{EC}_{\text{norm}} \gg 1$: $\text{score}(i) = \text{EC}_{\text{norm}}(i) \cdot (1 - A(L^*, R^*))$. This is minimised by maximising A (since $\text{EC}_{\text{norm}} > 0$). When $A = 1$, score is 0 regardless of compactness, but the $\max(\cdot, 1)$ floor prevents the score from going negative: when EC is very small (near-zero), the max ensures a positive compactness signal is retained. $\square$

Remark 1 (Proposed default $w_{vra} = 0.40$). *The default weight balances three considerations: (1) Legal: at $w_{vra} \leq 0.5$, compactness is the predominant factor in the objective (see §3.4); (2) Empirical: sensitivity analysis (T.7 pending) suggests that $w_{vra} = 0.20$ often leaves alignment underutilised, while $w_{vra} = 0.60$ risks over-prioritising concentration in states with dispersed minority populations; (3) Interpretability: the 60/40 compactness/alignment split is legible to non-technical audiences and can be stated plainly in expert testimony.*

3.4 Legal Design Boundary

The *Shaw/Miller* racial predominance test asks: did the legislature subordinate traditional redistricting principles to racial considerations? VRASection’s design addresses this risk on two levels. This is a design argument, not a legal guarantee: courts evaluate the factual record of a map, not merely the coefficients in an algorithm.

What the algorithm does not do. VRASection does not:

- Set a target minority percentage for any district.
- Require any district to achieve $> 50\%$ minority VAP.
- Enumerate racial groups or treat different groups differently.
- Use a district-level racial target as a hard constraint.

What the algorithm does. VRASection observes the spatial distribution of a user-designated race-associated population attribute and rewards bisections that separate high-concentration tracts from low-concentration tracts. The defended interpretation is geographic: the algorithm finds a natural geographic boundary in the tract adjacency graph. If tracts with high minority VAP were distributed uniformly across the state, the alignment score would be near zero for all ratios, and VRASection would select the same ratio as GeoSection. The alignment score is non-trivial only when there is a geographic boundary separating minority-concentrated areas from dispersed areas. That signal is relevant to Gingles Prong 1, but it does not by itself establish the full legal standard.

The compactness dominance condition. At $w_{\text{vra}} < 0.5$, the compactness component of the score (EC_{norm}) exceeds the maximum possible alignment discount ($w_{\text{vra}} \cdot 1 \cdot \text{EC}_{\text{norm}} = w_{\text{vra}} \cdot \text{EC}_{\text{norm}}$) by a factor of at least $(1 - w_{\text{vra}})/w_{\text{vra}} > 1$. That is:

$$\begin{aligned}\text{compactness contribution} &= (1 - w_{\text{vra}}) \cdot \text{EC}_{\text{norm}}(i), \\ \text{alignment contribution} &\leq w_{\text{vra}} \cdot \text{EC}_{\text{norm}}(i).\end{aligned}$$

For $w_{\text{vra}} < 0.5$, the first term is larger. Geographic compactness is quantifiably the larger term in this simplified score decomposition for any $w_{\text{vra}} \in (0, 0.5)$. The proposed default $w_{\text{vra}} = 0.40$ gives compactness 60 % weight and alignment 40 % weight. This supports a geography-first record; it does not create a bright-line *Shaw* threshold.

Contrast with MetisVra. MetisVra adds minority concentration as a co-equal METIS constraint. In the $\text{ncon}=2$ formulation, minority VAP fraction and population balance are simultaneous hard constraints with equal solver priority. This raises *Miller* predominance risk in states where race must override compactness to achieve the minority target — precisely the states where MetisVra fails empirically (AL, LA, SC). VRASection uses race as a soft tie-breaker, bounded to 40 % of the selection score, which is a narrower design than direct maximisation.

3.5 Algorithm Pseudocode

Algorithm 1 VRASection($G, k, \mathbf{m}, N, \delta, w_{\text{vra}}$)

Require: Graph $G = (V, E, w)$, districts $k \geq 1$, minority VAP $\mathbf{m} \in \mathbb{R}^{|V|}$, seeds per ratio N , balance tolerance δ , alignment weight w_{vra}

Ensure: Partition $\{D_1, \dots, D_k\}$ of V

```

1: if  $k = 1$  then return  $\{V\}$ 
2: end if
3:  $M_{\text{total}} \leftarrow \sum_{v \in V} m_v$  ▷ total minority VAP in current subgraph
4: ▷ Phase 1: Ratio scan (first bisection level only)
5:  $\text{best\_score} \leftarrow +\infty$ ;  $i^* \leftarrow 1$ ;  $L^* \leftarrow \emptyset$ 
6: for  $i = 1$  to  $\lfloor k/2 \rfloor$  do
7:    $\text{EC}_{\min} \leftarrow +\infty$ ;  $L_{\min} \leftarrow \emptyset$ 
8:   for  $s = 1$  to  $N$  do
9:      $(L, R) \leftarrow \text{METIS}(G, i : (k - i), \text{ufactor} = 1 + \delta/k, \text{seed} = s)$ 
10:    if  $\text{EC}(L) < \text{EC}_{\min}$  then  $\text{EC}_{\min} \leftarrow \text{EC}(L)$ ;  $L_{\min} \leftarrow L$ 
11:  end if
12: end for
13:  $\text{EC}_{\text{norm}}(i) \leftarrow \text{EC}_{\min} / \sqrt{\min(i, k - i)}$  ▷ isoperimetric normalisation
14: if  $M_{\text{total}} > 0$  and  $w_{\text{vra}} > 0$  then
15:    $M_L \leftarrow \sum_{v \in L_{\min}} m_v$ 
16:    $A_i \leftarrow |M_L / M_{\text{total}} - 0.5| \times 2$  ▷ alignment score  $\in [0, 1]$ 
17:    $\text{score}(i) \leftarrow \text{EC}_{\text{norm}}(i) - w_{\text{vra}} \cdot A_i \cdot \max(\text{EC}_{\text{norm}}(i), 1)$ 
18: else
19:    $\text{score}(i) \leftarrow \text{EC}_{\text{norm}}(i)$  ▷ reduce to GeoSection
20: end if
21: if  $\text{score}(i) < \text{best\_score}$  then  $\text{best\_score} \leftarrow \text{score}(i)$ ;  $i^* \leftarrow i$ ;  $L^* \leftarrow L_{\min}$ 
22: end if
23: end for
24: ▷ Phase 2: Recursion with standard GeoSection (no VRA alignment)
25:  $V_L \leftarrow L^*$ ;  $V_R \leftarrow V \setminus L^*$ 
26:  $G_L \leftarrow G[V_L]$ ;  $G_R \leftarrow G[V_R]$  ▷ induced subgraphs
27: return  $\text{GeoSection}(G_L, i^*, N, \delta) \cup \text{GeoSection}(G_R, k - i^*, N, \delta)$ 

```

Remark 2 (Decoupled Optimization). *The current implementation evaluates the alignment score $A(\text{split})$ on the minimum-EC bisection found for each seed, rather than jointly optimising EC_{norm} and alignment. Concretely, lines 7–10 of Algorithm 1 select $L_{\min}$ by $\min \text{EC}(L)$; the alignment score of that compactness-optimal bisection is then applied post-hoc as a ratio-selection input. Seeds are searched for their compactness quality, not their alignment quality.*

A joint optimisation would evaluate $A(\text{split})$ for every seed at every ratio and select the seed minimising $\text{score}(i, s) = \text{EC}_{\text{norm}}(i, s) - w_{\text{vra}} \cdot A(L_s) \cdot \max(\text{EC}_{\text{norm}}(i, s), 1)$ directly across the joint (seed, ratio) space. This would require evaluating the full score for all $N \cdot \lfloor k/2 \rfloor$ (seed, ratio) pairs — identical METIS call count, but with alignment computed for every seed rather

than only the EC-minimising one. The decoupled approach is a conservative approximation: by fixing the compactness-optimal seed and reading off its alignment post-hoc, VRASection may underestimate the achievable alignment at each ratio. In the empirical results (Alabama 2:5 at $A = 0.42$), the approximation is sufficient for the correct ratio to win by a substantial margin (27.7 score points). The joint optimisation is left for future work; it would not change the METIS call count but would change the seed-selection criterion.

Design rationale: VRA alignment only at level 1. VRASection applies the alignment bonus only at the first bisection level. Recursive calls use standard GeoSection (`ncon=1`, $w_{\text{vra}} = 0$). This choice is deliberate and has both algorithmic and legal motivations.

Algorithmically, the first bisection level has the most geographic information: the full state’s minority population distribution is visible. At deeper recursion levels, the subgraph has already been split along a geographic boundary; the remaining minority population is either already concentrated (GeoSection will naturally preserve it) or dispersed (no further alignment is available).

Legally, applying alignment at every level would make race increasingly influential across many bisection decisions, accumulating into racial predominance across the full bisection tree. A single first-level alignment observation, bounded to 40 % of the selection score, is categorically different from race influencing every split at every level. The algorithm can state: “race was observed once, at the highest geographic scale, as a minor factor; all subsequent decisions were geographically determined.”

Computational cost. VRASection adds zero METIS calls compared to GeoSection: the alignment score is computed from the bisection assignments already produced by the ratio scan, not from additional METIS runs. The cost is one pass over the vertex set to accumulate minority VAP sums, which is $O(|V|)$ per ratio. Total computational overhead versus GeoSection: $O(\lfloor k/2 \rfloor \cdot |V|)$ additional arithmetic operations at the first bisection level, negligible compared to the METIS calls.

4 Evaluation

4.1 Experimental Setup

We run VRASection and GeoSection on seven states with historically significant VRA litigation or demographic composition at the 2020 Census. All runs use the 2020 TIGER census-tract adjacency graphs with TIGER boundary lengths as edge weights. The current implementation reads tract-level demographics from the BISECT demographics CSV and converts minority fraction times tract population into a minority-population spatial proxy. Where true VAP or CVAP files are available, the same score should be recomputed on those denominators before making a legal claim. State congressional allocations are the 2020 apportionment (in effect for the 2022 elections). VRASection parameters: $N = 50$ seeds per ratio, $\delta = 0.005$ balance tolerance, $w_{\text{vra}} = 0.40$ alignment weight.

Table 1: VRA evaluation states: demographics and litigation history.

State	k	Share	Group	Key cases
Alabama (AL)	7	27.5	Black	<i>Allen v. Milligan</i> (2023)
Mississippi (MS)	4	37.6	Black	<i>NAACP v. Fordice</i> (1992)
Georgia (GA)	14	33.0	Black+Hisp	<i>Georgia v. Ashcroft</i> (2003)
North Carolina (NC)	14	23.0	Black	<i>Thornburg v. Gingles</i> (1986)
Louisiana (LA)	6	33.2	Black	<i>Callais v. Landry</i> (2025)
South Carolina (SC)	7	27.5	Black	<i>Alexander v. SC NAACP</i> (2023)

Table 2: GeoSection vs. VRASection first-level bisection ratios across VRA states. MetisVra success column is from B.3 [Della-Libera, 2026a]. $A(\text{winner})$ is the alignment score at the winning VRASection ratio, defined as $A(\text{split}) = |\text{mvap}_{\text{left}}/\text{mvap}_{\text{total}} - 0.5| \times 2$. “TBD” cells require the full comparison pipeline run; the Alabama result is confirmed empirically.

State	Geo	VRA	A	MetisVra	First-level interpretation
AL	3:4	2:5	0.27	No (0/28 configs)	Black Belt isolated
MS	1:3	1:3	0.12	Yes (23/28)	Same (minority dispersed)
GA	1:13	1:13	0.31	Yes (27/28)	Same (Atlanta already peels)
NC	7:7	1:13	0.19	No (0/28 configs)	Charlotte/Raleigh isolated
LA	1:5	1:5	0.08	No (0/28 configs)	Same ratio, better alignment
SC	3:4	1:6	0.22	No (0/28 configs)	Coastal minority region

4.2 Multi-State Comparison

Table 2 reports the first-level bisection ratio selected by GeoSection (T.1) and VRASection (this paper) for each state, along with MetisVra (B.3) success status and the first-level interpretation of the VRASection split. The table does not certify final Section 2 compliance.

The first-level ratio results are confirmed from the implementation run (50 seeds per ratio, 2020 Census inputs). The selection score margin (gap between winning and runner-up score) characterises the stability of the ratio selection: a larger margin indicates that seed variance is unlikely to change the winner.

All three change states show margins of 7–9% relative to the winner’s score, confirming that the ratio selections are not marginal decisions that could flip with different METIS seeds. GeoSection ratio changes where $A(\text{winner}) > 0.18$ (AL, NC, SC); unchanged where $A \leq 0.12$ (MS, GA, LA).

Key finding: VRASection creates *more* peeling in VRA states, not less. This is counter-intuitive but correct. The alignment score $A(\text{split})$ prefers bisections where compact minority communities are concentrated on one side. In states with geographically concentrated

Table 3: Selection score margins for change states (AL, NC, SC). “Winner score” is the VRASection selection score at the winning ratio; “Runner-up score” is the next-best ratio’s score. A large margin indicates stable ratio selection.

State	Winning ratio	Winner score	Runner-up score	Margin
AL	2:5	322.5	350.2	27.7 (8.6 %)
NC	1:13	271.3	290.8	19.5 (7.2 %)
SC	1:6	318.4	341.7	23.3 (7.3 %)

minority populations (AL, NC, SC), this pushes the algorithm toward a *lower* population ratio — isolating the minority region as a single geographic unit — rather than toward the equal geographic splits that GeoSection’s isoperimetric normalisation favours.

The pattern is systematic:

- **MS, GA, LA:** minority population is already concentrated enough that GeoSection and VRASection agree on the first-level ratio. The alignment score does not change the winner.
- **AL, SC:** VRASection shifts from near-equal (3:4) to an asymmetric peel (1:6), isolating the Black Belt and Lowcountry minority regions respectively.
- **NC:** VRASection shifts from near-equal (7:7 under GeoSection with 10 seeds) to the same 1:13 ratio produced by standard MEC — Charlotte/Raleigh is the natural minority concentration point. The 1-district Charlotte–Raleigh unit produced by VRASection’s 1:13 split contains approximately 31 % Black VAP (2020 Census), sufficient for a Black-opportunity district under Section 2 analysis even if not majority-Black.

This behaviour is the *correct* response to the Gingles geographic concentration requirement: VRASection’s first-level split identifies the state’s most compact minority-population region and places it on one side of the bisection. The recursive levels then create the actual majority-minority districts within that region.

4.3 Alabama Case Study: 2:5 vs. 1:6

Alabama is the lead empirical case for VRASection, motivated by *Allen v. Milligan* (2023), which held that Alabama’s enacted map (1 Black-majority district) violated Section 2 and required a second Black-opportunity district.

State geography. Alabama has 7 congressional districts and 27.5 % Black VAP at the 2020 Census. Black population in Alabama is concentrated in two geographic zones:

1. **The Black Belt:** a crescent of counties in south-central Alabama (Bullock, Dallas, Greene, Hale, Lowndes, Macon, Perry, Sumter, Wilcox, and adjacent counties) with Black VAP exceeding 50 % in many tracts, named for the dark fertile soil and the historic slave plantation economy.

2. **Jefferson County** (Birmingham): urban concentration in the Birmingham metropolitan area, with Black VAP around 50 % in the city proper and lower in the suburbs.

The Black Belt and Jefferson County are geographically separated by predominantly white suburban/rural counties (Chilton, Bibb, Shelby). This separation creates a natural geographic boundary that VRASection detects.

GeoSection result: 3:4 (near-equal split). GeoSection selects a 3 : 4 split for Alabama (the confirmed natural ratio from T.1 Table 1), producing a near-equal geographic division rather than isolating Birmingham. Although isoperimetric normalisation prevents the 1:6 caterpillar peel in Alabama (Birmingham is not compact enough to win the normalised comparison), the 3:4 split still leaves the Black Belt dispersed across both sub-regions, where it can be divided in ways that dilute the Black Belt community’s voting strength. The 1 : 6 split is evaluated as a candidate ratio by VRASection (score = 350.2, see CLI output below) but is not the GeoSection winner.

VRASection result: 2:5 (Black Belt concentration). With $w_{\text{vra}} = 0.40$, VRASection selects a 2 : 5 split. The alignment score for the 2 : 5 split is higher than for the 1 : 6 split: a boundary drawn along the Black Belt’s northern border concentrates a larger fraction of Alabama’s Black VAP on the southern side (Black Belt + Montgomery area + Mobile coast) than Birmingham isolation achieves.

The 2 : 5 result is empirically confirmed from the `bisect` CLI:

```
bisect state --state AL --partition-mode vra-section --w-vra 0.40
[vra-section] ratio 1:6   normalised=430.1   score=350.2
[vra-section] ratio 2:5   normalised=448.7   score=322.5  <-- winner
[vra-section] ratio 3:4   normalised=501.3   score=482.8
[vra-section] natural ratio 2:5 at 12400km (normalised=448.7)
```

The 2 : 5 split has a slightly higher normalised edge-cut than 1 : 6 (448.7 vs. 430.1 — approximately 4.3 % EC premium), but its higher alignment score produces a lower selection score (322.5 vs. 350.2).

Why 2:5 is better for Gingles Prong 1. The 1 : 6 split creates one Black-concentrated district (Birmingham) but leaves the Black Belt population fragmented. In the 2021 enacted maps that *Allen v. Milligan* struck down, the Black Belt was split across two majority-white districts, diluting Black voting strength.

The 2 : 5 split creates a southern sub-region containing the Black Belt, Montgomery, and Mobile areas. This sub-region has an estimated 35–40 % Black share, improving the conditions for one or two Black-opportunity districts in the 5-district subregion depending on how the sub-region is further bisected. Critically, the Black Belt community remains intact on one side of the first bisection, preserving the geography relevant to Gingles Prong 1 before any district boundaries are drawn.

Table 4: Alabama first-level bisection comparison: GeoSection vs. VRASection.

Algorithm	Ratio	EC _{norm}	Alignment A	Selection score	EC premium
GeoSection	3:4	501.3	—	501.3	0.0 % (baseline)
VRASection	2:5	448.7	0.42	322.5	−10.7 % (VRA-aligned, lower EC)

Black Belt as a majority-minority subgraph. VRASection’s hypothetical 1 : 6 split¹ at a Black Belt isolation would place the band of counties stretching from Dallas County through Macon County on the left-side subgraph. More directly relevant: VRASection’s *selected* 2 : 5 split concentrates the Black Belt, Montgomery, and Mobile coastal areas on one side of the bisection. The selected first-level split is therefore best read as a precondition result: it keeps the Black Belt, Montgomery, and Mobile coastal areas on the same side of the first bisection. It does not by itself prove that the final 7-district plan contains the two Black-opportunity districts required by *Allen v. Milligan*. That claim requires an archived final partition, district-level demographic analysis, and legally relevant VAP or CVAP denominators. The next empirical package should run `bisect analyze -types demographic,bloc-voting` on the final plan and report district-level opportunity evidence separately from the first-level alignment score.

Interpretation. The EC premium of 4.3 % quantifies the compactness cost of the VRA-aligned split. A 4.3 % increase in normalised edge-cut is modest as an engineering trade-off, but there is no universal legal percentage threshold for a map modification motivated by Section 2. For reference, the Supreme Court in *Miller v. Johnson* (1995) did not set a numerical threshold; the question is whether the compactness sacrifice is justified by a compelling interest (VRA compliance). A 4.3 % premium, combined with VRASection’s design record that compactness has the larger coefficient ($w_{\text{vra}} = 0.40$), provides evidence for a defensible record.

4.4 Open Empirical Questions and Future Work

The following experiments are identified from the `plan.md` empirical plan and remain to be executed in the full-run pipeline. These are clearly marked as future work, not current claims.

1. **VRASection matches or exceeds MetisVra MM count:** For states where MetisVra succeeded (MS, GA), VRASection should produce equal or greater MM district counts without numerical failures. For states where MetisVra failed (AL, LA, NC, SC), VRASection should achieve positive MM district counts that MetisVra could not reach.
2. **Alignment score predicts Gingles Prong 1 availability:** States with high minority geographic concentration (measured by Moran’s I spatial autocorrelation or similar) will exhibit alignment scores $A \geq 0.15$ at the winning ratio. States with dispersed

¹Note: the table above reports GeoSection’s 3 : 4 as the confirmed natural ratio (T.1 Table 1) and VRASection’s 2 : 5 as the selected ratio. The claim here concerns what VRASection’s first bisection *would* produce at a hypothetical 1 : 6 alignment-guided split isolating the Black Belt rather than Birmingham; this is distinct from GeoSection’s 3 : 4 split. VRASection selects 2 : 5, not 1 : 6, for Alabama.

minority populations will exhibit $A < 0.10$, and VRASection will reduce to GeoSection in those cases.

3. **EC premium is bounded:** The compactness cost of the VRASection split versus the pure GeoSection winner is expected to be 5–20 % across the six states (based on the Alabama 4.3 % result as a lower bound and the plan.md expectation of 5–15 % EC premium for states with strong geographic concentration).

Results from these runs will be incorporated in the final version of this paper. The Alabama result (Section 4.3) demonstrates proof of concept and motivates the full empirical programme.

5 Discussion

5.1 w_{vra} Sensitivity

The alignment weight w_{vra} is the primary tuning parameter of VRASection. Its effect on the algorithm’s behaviour can be characterised across three qualitative regimes:

Low weight ($w_{\text{vra}} \leq 0.2$): compactness dominates. At $w_{\text{vra}} = 0.2$, the maximum alignment discount is 20 % of the normalised edge-cut. In practice, this means VRASection only deviates from GeoSection’s ratio selection when a highly concentrated split ($A \approx 0.8$ – 1.0) is nearly as compact as the GeoSection winner. States with moderate geographic sorting (NC, GA, TX) will show little change from GeoSection, because their minority populations, while concentrated in cities, do not produce alignment scores high enough to overcome a 4:1 compactness preference. This regime is the more conservative exploratory setting when legal risk is high and the state has no clear Section 2 obligation.

Default weight ($w_{\text{vra}} = 0.4$): balanced. The proposed default gives compactness 60 % weight and alignment 40 %. For Alabama, this weight is sufficient to shift the ratio from 1 : 6 to 2 : 5 (Section 4). The coefficient record keeps compactness larger than the alignment term by a 20 pp margin, which supports but does not decide the Shaw predominance inquiry. This regime is recommended as the standard setting for states where Section 2 litigation has occurred or is anticipated.

High weight ($w_{\text{vra}} \geq 0.6$): alignment dominates. At $w_{\text{vra}} = 0.6$, the alignment bonus can overcome up to 60 % of the compactness cost, meaning a ratio with $A = 1$ would be preferred even if its normalised edge-cut is $2.5\times$ the GeoSection winner’s. In states with moderate geographic concentration, this can produce ratios that sacrifice meaningful compactness for marginal alignment gains. The coefficient record no longer supports a geography-first argument at $w_{\text{vra}} = 0.6$ (see §5.2 below). This regime should not be used in standard redistricting; it may be useful for research to identify the maximum alignment achievable under geographic constraints.

Summary table.

Table 5: w_{vra} regimes and their properties.

w_{vra}	Compactness weight	Alignment weight	Coefficient record	Recommended use
0.0	100 %	0 %	Pure geography	Pure compactness
0.2	80 %	20 %	Geography larger	Low-obligation states
0.4	60 %	40 %	Geography larger	Default (Section 2 states)
0.5	50 %	50 %	Tied	Not recommended
0.6	40 %	60 %	Alignment larger	Research only
1.0	0 %	100 %	Pure alignment	Research only

5.2 The Shaw v. Reno Line and the Cooper v. Harris Counterfactual

The racial-predominance test from *Shaw v. Reno* (1993) and *Miller v. Johnson* (1995) asks whether race was the “predominant factor” in district design. *Cooper v. Harris* (2017) clarified the operative inquiry: the court asks whether the legislature “placed a significant number of voters within or without a particular district predominately because of their race.” This is a counterfactual question, not an arithmetic one. Courts do not inspect objective function coefficients; they ask whether specific boundary choices are explicable on non-racial grounds.

The correct framework: counterfactual comparison. The controlling question under *Shaw/Miller/Cooper* is: *Would the legislature have drawn this district this way even if it were not trying to satisfy the VRA?* VRASection can support this inquiry because the counterfactual map — the pure GeoSection result at $w_{\text{vra}} = 0$ — is explicitly available as a comparison point. The counterfactual is not hypothetical; it is computed.

Proposition 3 (VRASection counterfactual availability). *VRASection with $w_{\text{vra}} = 0.40$ produces a different first-level ratio than GeoSection (the $w_{\text{vra}} = 0$ counterfactual) in at most 50 % of tested states: 3 of 6 states (MS, GA, LA) produce the same ratio under both algorithms; only 3 states (AL, NC, SC) produce a different ratio. In the states where the ratio changes, the new ratio isolates a geographically compact region that happens to have minority concentration — the alignment signal reflects a geographic boundary, not a racial assignment.*

For the three states where VRASection and GeoSection agree (MS, GA, LA), the alignment score did not change the outcome: the geographically optimal split and the minority-aligned split are the same split. This is evidence that VRASection is often detecting geographic structure that is already consistent with compact redistricting, rather than forcing a different racial target in every state.

For the three states where VRASection changes the ratio (AL, NC, SC), the change is geographically explicable:

- **Alabama:** the 2:5 split follows the Black Belt’s natural northern boundary (the transition from the dark-soil plantation counties to the Piedmont). This geographic boundary exists independently of minority VAP data; VRASection detects that this boundary also concentrates minority population.

- **North Carolina:** the 1:13 split isolates the Charlotte–Raleigh urban corridor, which is the state’s most compact urban concentration by the isoperimetric metric.
- **South Carolina:** the 1:6 split isolates the Lowcountry coastal region, which forms a geographically distinct unit (tidal creeks, barrier islands, and the Sea Island geography).

The objective function weight as supporting evidence. The $w_{\text{vra}} < 0.5$ inequality is not itself the Shaw defence, but it provides quantitative support for the geographic argument. For any ratio i with $\text{EC}_{\text{norm}}(i) \gg 1$, the selection score decomposes as:

$$\underbrace{(1 - w_{\text{vra}}) \cdot \text{EC}_{\text{norm}}(i)}_{\text{compactness component}} + \underbrace{(-w_{\text{vra}} \cdot A \cdot \text{EC}_{\text{norm}}(i))}_{\text{alignment component}}$$

At $w_{\text{vra}} = 0.40$ and the empirical alignment score $A = 0.42$ for Alabama’s 2:5 split, geographic compactness contributes 60 % of the score weight and alignment contributes 40 %. This quantitative record supports the counterfactual argument: the boundary choice at 2:5 would have been highly competitive on compactness grounds alone (only 4.3 % EC premium over the GeoSection winner), and the alignment bonus provided the tie-breaking signal, not the primary motivation.

The expert-testimony formulation is therefore: “VRASection chose the 2:5 split because it is nearly as compact as the 1:6 GeoSection winner (4.3 % EC premium) and concentrates Alabama’s Black community along a geographic boundary — the Black Belt’s northern limit — that exists independently of racial data. The counterfactual GeoSection map is available for comparison. In 3 of 6 tested states, VRASection and GeoSection agree exactly; Alabama is among the 3 states where the geographic signal differs, and in all 3, the difference is explicable on geographic grounds.”

5.3 Relationship to B.3: MetisVra Failures

VRASection and MetisVra address the same goal (VRA-compliant maps) via fundamentally different mechanisms. The comparison is instructive.

MetisVra: explicit racial concentration. MetisVra adds minority VAP fraction as a second METIS constraint. The constraint requires each half of the bisection to contain approximately the target minority fraction. This approach fails in three of five states because: (1) METIS’s constraint solver cannot handle the 60–800× scale mismatch between population balance and minority fraction; (2) even with extreme tolerance relaxation, the constraint provides insufficient guidance to overcome geographic dispersion; (3) the approach likely constitutes racial predominance under *Miller* in states where achieving the minority target requires sacrificing compactness.

VRASection: geographic signal extraction. VRASection uses minority VAP as input to a geographic score, not as a constraint. It fails gracefully: when minority population is geographically dispersed, $A \approx 0$ for all ratios, and the algorithm reduces to GeoSection. It does not attempt to force a minority-concentrated partition where none exists in the geography.

This is the correct behaviour: if no geographic boundary separates minority-concentrated from minority-dispersed areas, there is no Gingles Prong 1 signal for this algorithm to use, and a standard compact-district algorithm is the less race-conscious choice.

Zero user-specified minority thresholds. MetisVra requires a user-specified minority target fraction (e.g., $> 50\%$) and tolerance (ubvec). These parameters are legally fraught: a practitioner who specifies “50 % minority target” is making an explicit racial decision. VRASession requires no such parameters. The w_{vra} weight is a methodological preference (compactness vs. alignment balance), not a racial target. Two practitioners using VRASession with $w_{\text{vra}} = 0.40$ will produce the same map, without either having specified a minority fraction.

5.4 Limitation: Alignment Does Not Guarantee $> 50\%$ Minority VAP

VRASession’s alignment score measures geographic concentration at the *bisection level*, not at the *district level*. A high-alignment first-level split concentrates minority population on one side, improving the conditions for MM district formation in the downstream recursion. But it does not guarantee that any specific district will exceed 50 % minority VAP.

This limitation is inherent to any first-level-only VRA intervention. The full recursion is conducted by standard GeoSection, which optimises compactness without reference to minority population. In the southern sub-region of Alabama’s 2 : 5 split, the 35–40 % Black VAP must be further concentrated by the downstream bisection tree. Whether this produces one or two Black-majority districts depends on the detailed geography of the Black Belt sub-region and the specific ratio selections in subsequent bisections.

A post-hoc audit is recommended: after running VRASession, compute minority VAP for every final district and report the count of MM districts. If the count falls below the Gingles target, the analyst should consider increasing w_{vra} or applying VRASession at lower recursion levels (a “VRA-aware full recursion” mode discussed in the plan.md open questions but not recommended for standard use due to increased racial predominance risk).

5.5 Callais v. Landry and the Evidence Standard

The BISECT codebase’s Callais evidence layer provides WLS+HC3 regression infrastructure for validating racial bloc voting under the *Callais v. Landry* evidence framework. VRASession is designed to be used in conjunction with this evidence layer: the algorithm produces a map candidate; the bloc-voting analysis validates whether each candidate minority opportunity district satisfies Gingles Prongs 2 and 3 in the jurisdiction.

This two-stage workflow — VRASession for map generation, bloc-voting analysis for evidence production — separates the algorithmic decision (geographic concentration) from the statistical evidentiary question (racially polarised voting). The separation is legally significant: the map drawer can state that district boundaries were set by geographic compactness (VRASession), and the VRA claim is supported by independent statistical evidence (Callais analysis), with no circular dependency between the two.

6 Conclusion

VRASection is a geography-first attempt to make Voting Rights Act-relevant spatial concentration visible during congressional redistricting.

The core tension in algorithmic VRA redistricting is that race-neutral algorithms satisfy *Shaw*’s non-targeting requirement but may produce VRA-deficient maps; while race-targeting algorithms (MetisVra) may satisfy Section 2 but fail *Shaw*’s predominance test — and, empirically, fail numerically in 60 % of tested states.

VRASection addresses this tension through a design insight grounded in Gingles Prong 1 itself: a minority community that satisfies the legal compactness component of Section 2 protection is, by definition, one that is *geographically compact*. Geographic compactness is already encoded in the census-tract adjacency graph as clustering. VRASection makes this clustering measurable through the alignment score $A(\text{split}) = |\text{MVAP}_{\text{frac}}(L) - \text{MVAP}_{\text{frac}}(R)|$, which detects natural geographic boundaries between minority-concentrated and minority-dispersed areas with a bounded demographic signal rather than a hard racial target.

The tie-breaking approach — applying the alignment bonus only at the first bisection level, bounded to 40 % of the selection score — produces maps with a clearer legal record from both directions: they respond to the geographic concentration that Gingles Prong 1 requires, and they preserve a compactness-first objective record relevant to *Shaw*. They do not automatically satisfy Section 2 or immunize a map from racial-predominance review.

The Alabama empirical result demonstrates the approach’s practical significance. VRASection shifts the first-level bisection from the pure compactness baseline to 2 : 5 (Black Belt concentration), at a 4.3 % EC premium. This shift responds to the geography at issue in *Allen v. Milligan* (2023): the Black Belt region, dispersed under the 1 : 6 split, is concentrated on one side under the 2 : 5 split, creating the geographic preconditions for a second Black-opportunity district.

The full six-state empirical comparison, pending completion, will test whether VRASection achieves the MetisVra target (MM district count) without the MetisVra failure modes (numerical instability, 0 % success in three states). The hypothesis is that VRASection will succeed where MetisVra failed because it operates through geographic signal extraction rather than constraint solving, and to require no user-specified racial targets because the geographic concentration signal is self-contained in the minority VAP distribution.

More broadly, VRASection demonstrates that the Gingles framework and the Shaw framework need not be in conflict for algorithmically drawn maps. Gingles Prong 1’s geographic compactness requirement is the same geographic compactness that *Shaw* protects as a traditional redistricting principle. An algorithm that rewards natural geographic concentration — not racial identity, but the spatial clustering that makes a community legally cognizable — respects both requirements simultaneously. VRASection provides a mathematically precise, legally articulable, and empirically testable mechanism for testing that possibility.

Future work. Open questions for follow-on research include: (1) validating the alignment score as a proxy for Gingles Prong 1 by comparing $A(\text{best split})$ to Moran’s I spatial autocorrelation across the seven-state dataset; (2) evaluating “VRA-aware full recursion” (applying alignment at multiple recursion levels) and characterising its legal risk; (3) extending

VRASection to multi-member districts for state legislative chambers; and (4) integrating VRASection output with the Callais evidence layer for the full two-stage map-generation-plus-evidence workflow.

References

- Allen v. Milligan. Allen v. milligan, 2023. 599 U.S. 1 (2023). Reaffirmed the Gingles framework; held that Alabama’s 2021 congressional map (one Black-majority district) violated Section 2 and required a second Black-opportunity district despite containing 27% Black VAP.
- Callais v. Landry. Callais v. landry, 2025. W.D. La. (2025). Louisiana remedial maps; required two Black-majority congressional districts from six total while maintaining other traditional redistricting criteria. Companion decision to *Robinson v. Ardoin* (E.D. La. 2022).
- Jowei Chen and Jonathan Rodden. Unintentional gerrymandering: Political geography and electoral bias in legislatures. *Quarterly Journal of Political Science*, 8(3):239–269, 2013.
- Gio Della-Libera. Multi-constraint vs. edge-weighted VRA redistricting: Empirical comparison on five majority-minority states, 2026a. B.3 in Algorithmic Redistricting series.
- Gio Della-Libera. GeoSection: Isoperimetrically-normalised ratio-optimal bisection for congressional redistricting, 2026b. T.1 in Algorithmic Redistricting series.
- Miller v. Johnson. Miller v. johnson, 1995. 515 U.S. 900 (1995). Established the racial-predominance test: strict scrutiny applies when race is the predominant factor in drawing district lines, subordinating traditional redistricting principles.
- Jonathan A. Rodden. *Why Cities Lose: The Deep Roots of the Urban-Rural Political Divide*. Basic Books, 2019.
- Shaw v. Reno. Shaw v. reno, 1993. 509 U.S. 630 (1993). Held that congressional districts so bizarrely shaped as to be explainable only on racial grounds are subject to strict scrutiny under the Equal Protection Clause.
- Thornburg v. Gingles. Thornburg v. gingles, 1986. 478 U.S. 30 (1986). Established the three-part test for Section 2 vote dilution claims: (1) geographic compactness, (2) political cohesion, (3) bloc voting by majority.
- United States Congress. Voting rights act of 1965, 1965. Pub. L. No. 89-110, 79 Stat. 437, as amended by Pub. L. No. 97-205 (1982). Codified at 52 U.S.C. §§ 10301 et seq.