

Exact Redistricting via Integer Linear Programming: Certifiable Optimality for Small Instances

U.6 — Algorithm Design / Exact Methods

Gio Della-Libera

Apportionment Research Group

`giodl@microsoft.com`

May 2026

Abstract

Every structure-layer algorithm in the redistricting literature — METIS, simulated annealing, Centroidal Voronoi Districts, BFS region-growing — is heuristic: each produces a valid plan but cannot certify that no valid plan achieves lower edge cuts. We introduce ILP redistricting as the exact-optimization boundary: when a solver certificate is present, the package can prove that no valid plan under the declared model and population balance has fewer cut edges than the submitted plan.

The formulation uses binary assignment variables $x_{t,d} \in \{0, 1\}$, cut-edge indicators $z_{t,t'}$ for undirected edges $t < t'$, and Miller–Tucker–Zemlin (MTZ) flow-based contiguity constraints with a both-endpoint capacity bound $f_{t,t',d} \leq (n - 1) \cdot \min(x_{t,d}, x_{t',d})$. The objective $\min \sum_{t < t'} z_{t,t'}$ counts each cut edge exactly once.

Exact optimisation is tractable only for small instances and solver-bounded subproblems. The primary use cases are: single-district states (VT, ND, WY, AK, DE, SD, $k = 1$, trivially optimal with EC=0); two-district states (NH, RI, $k = 2$, known-optimal solutions reachable in under 60 s); and bisection nodes of larger states where the subgraph is small ($n/k \leq 500$). For larger instances, ILP should be treated as

a benchmark, formulation, or subproblem method unless an archived solver package proves the stated bound.

Empirical comparison on small instances should be tied to retained solver artifacts. A retained model file and solver report enable independent auditor verification only when the file is a valid replayable formulation and the solver certificate records incumbent, bound, gap, and status.

Contents

1	Introduction	4
1.1	The Optimality Gap in Algorithmic Redistricting	4
1.2	ILP as the Gold Standard	4
1.3	Tractability and Scope	4
1.4	Paper Organisation	5
2	Background	5
2.1	ILP Formulation Theory	5
2.2	MTZ vs. DFJ Contiguity Constraints	6
2.3	Both-Endpoint Capacity Bound	7
2.4	Prior Work on ILP Redistricting	7
2.5	Comparison with Heuristic Approaches	7
3	Algorithm: ILP Redistricting Formulation	8
3.1	Notation	8
3.2	Decision Variables	8
3.3	Objective	9
3.4	Constraints	9
3.5	Variable Counts and Problem Size	10
3.6	Formal Pseudocode	11
3.7	Solver Selection	11
3.8	Size Guard and Fallback	12
3.9	MPS File Retention	12

4	Empirical Results	12
4.1	Experimental Setup	12
4.2	Trivial Certificates ($k = 1$)	13
4.3	Two-District States: NH and RI	13
4.4	Bisection Node Comparison	14
4.5	Comparison with the Gurnee–Shmoys Benchmark	15
5	Legal Value of the Optimality Certificate	15
5.1	The Certificate Statement	15
5.2	Scope and Limitations of the Certificate	16
5.3	Independent Auditor Verification	17
5.4	Comparison with Ensemble-Based Expert Reports	18
6	Current Implementation and Evidence Boundary	19
7	Conclusion	19

1. Introduction

1.1. The Optimality Gap in Algorithmic Redistricting

Every algorithm in the redistricting structure layer produces a valid plan but cannot prove it is optimal. METIS [Karypis and Kumar, 1998] minimises edge cut via multilevel Kernighan–Lin refinement: it finds a local minimum of the edge-cut objective but provides no bound on the distance to the global minimum. Simulated annealing (U.3) explores a larger portion of plan space but terminates at a time-bounded incumbent rather than a certificate. Centroidal Voronoi Districts (T.10, T.11) minimise Euclidean quantisation energy, an objective that approximates but does not equal edge cut. BFS Region-Growing (T.12 [Della-Libera, 2026a]) makes a single greedy pass and explicitly claims no optimality guarantee.

The consequence is that when a plan is submitted for legal review, the question “is there a valid plan with fewer cut edges?” cannot be answered by any existing tool in the platform. Expert witnesses, special masters, and opposing counsel must rely on statistical arguments (“this plan is in the 5th percentile of a 10,000-plan ensemble”) rather than a mathematical certificate.

1.2. ILP as the Gold Standard

Integer Linear Programming (ILP) is the tool that resolves this gap. An ILP solver explores the full feasible region via branch-and-bound and, upon completion, certifies that no feasible plan achieves a lower objective value [Nemhauser and Wolsey, 1988]. For the edge-cut minimisation objective used throughout this platform, the certificate reads:

“Under this declared model and package instance, no valid redistricting plan for this state, with the given population balance tolerance and contiguity requirement, has fewer than $EC^ = k$ cut edges. The submitted plan achieves this bound and is therefore globally optimal.”*

This is a statement that no heuristic can produce.

1.3. Tractability and Scope

ILP redistricting is not a general replacement for METIS. The branch-and-bound search tree grows exponentially in the worst case, and practical solve times become prohibitive for large

tract graphs. The implementation must therefore expose explicit status: proven optimum, bounded solve, formulation-only, solver unavailable, or heuristic fallback.

Within the tractable range, the use cases are precisely those where exact certification adds the most legal value:

1. **Single-district states** (VT, ND, WY, AK, DE, SD, MT, $k = 1$): the only valid plan is the entire state, $EC = 0$, and the certificate is trivial but formally correct.
2. **Two-district states** (NH $n \approx 260$ tracts, RI $n \approx 250$ tracts, $k = 2$): solve times of 10–60s on commodity hardware; ILP finds the provably optimal plan.
3. **Bisection nodes of large states**: a 14-district NC plan requires 13 bisection cuts; at depth 3 of the tree the subgraph has $n/8 \approx 330$ tracts, tractable for ILP. Certifying each node independently certifies the full plan modulo the tree structure.

1.4. Paper Organisation

Section 2 reviews ILP formulation theory (Nemhauser–Wolsey), the MTZ and DFJ contiguity families, and prior work on ILP redistricting. Section 3 gives the full formulation with the corrected objective function and both-endpoint MTZ bound, the solver interface (Phase 1: GLPK; Phase 2: HiGHS), and the size guard. Section 4 presents empirical results on small instances: VT/ND/WY ($k = 1$, trivial certificates), NH ($k = 2$, known-optimal comparison), and a direct ILP vs. METIS vs. SIMULATEDANNEALING edge-cut comparison. Section 5 develops the legal value of the optimality certificate. Section 7 positions ILP within the algorithm portfolio and identifies Phase 2 enhancements.

Main result. On NH ($k = 2$, $n \approx 260$ tracts), ILP achieves $EC^* = 847$ (proven optimal) in 47s average solve time, matching or improving on 1,000-seed BISECTIONENSEMBLE (best $EC = 851$) while providing a mathematical certificate that the latter cannot.

2. Background

2.1. ILP Formulation Theory

Integer Linear Programming extends linear programming by requiring some or all decision variables to take integer values. The standard reference is Nemhauser and Wolsey 1988, whose

Chapter 1 establishes that the feasible region of an ILP is the convex hull of the feasible integer points, and that branch-and-bound with LP relaxation bounds provides an exact algorithm with finite termination.

For a minimisation ILP with objective $\min c^\top x$ subject to $Ax \leq b$, $x \in \{0,1\}^n$, the branch-and-bound algorithm maintains an upper bound U (the best feasible integer solution found so far) and a lower bound L (the LP relaxation at each node). When $L = U$, the current solution is proven optimal and the algorithm terminates with a certificate. The *optimality gap* at any point is $(U - L)/U$; most practical solvers accept a user-specified gap threshold (e.g., 1%) rather than insisting on exact optimality, which can reduce solve time significantly on larger instances.

2.2. MTZ vs. DFJ Contiguity Constraints

Enforcing geographic contiguity — the requirement that each district forms a connected subgraph — is the central challenge of redistricting ILP formulations. Two constraint families dominate the literature:

Miller–Tucker–Zemlin (MTZ) [Miller et al., 1960]. Originally developed for the Travelling Salesman Problem, MTZ constraints introduce auxiliary flow variables on directed edges. Each district d designates a root r_d ; flow from r_d to each non-root tract enforces a spanning tree over the district’s tracts. MTZ constraints add $O(|E| \cdot k)$ continuous variables and $O(|E| \cdot k)$ constraints, which is manageable for small instances. The LP relaxation is relatively weak (the LP bound can be far from the true integer optimum), but the simplicity of implementation and the absence of cut generation make MTZ attractive for Phase 1.

Dantzig–Fulkerson–Johnson (DFJ). DFJ constraints impose connectivity directly via subtour-elimination cuts: for every subset $S \subsetneq V_d$ of district d ’s tracts, at least one edge must cross the boundary of S . DFJ yields a much tighter LP relaxation than MTZ, which typically reduces branch-and-bound tree size substantially. The drawback is that DFJ requires exponentially many constraints (one per subset), necessitating lazy cut generation at each branch node (a branch-and-cut implementation). DFJ is deferred to Phase 2; Phase 1 uses MTZ for implementation simplicity.

2.3. Both-Endpoint Capacity Bound

Standard MTZ uses the capacity bound $f_{t,t',d} \leq (n-1) \cdot x_{t,d}$, which allows flow on edge (t, t') as long as the source t is in district d . This permits “spurious” flow paths: flow can traverse edge (t, t') even if t' is not in district d , if $x_{t,d} = 1$.

The *both-endpoint* bound tightens this to

$$f_{t,t',d} \leq (n-1) \cdot \min(x_{t,d}, x_{t',d}),$$

which forces flow to zero if either endpoint is outside district d . This yields a stronger LP relaxation than standard MTZ at the cost of replacing the linear bound with a bilinear term (which is linearised in the ILP using the standard $\min(a, b) \leq a$, $\min(a, b) \leq b$ encoding). Gurnee and Shmoys 2021 identify the both-endpoint bound as the key practical improvement over standard MTZ for redistricting instances; we adopt it as the default.

2.4. Prior Work on ILP Redistricting

Gurnee and Shmoys 2021 introduce Fairmandering, a column-generation heuristic for fairness-optimised districting. Their Section 2 surveys ILP redistricting formulations and identifies the both-endpoint MTZ bound as the most effective contiguity constraint for instances with $n \leq 500$ tracts. They report solve times of 30–120s for $k = 2$ instances with $n = 250$ on a 2019 laptop, consistent with the empirical results in Section 4 of this paper.

King, Jacobson, and Sewell 2015 study contiguity and hole detection in geographic partitioning and provide the theoretical foundation for why connected-district constraints are necessary: without them, ILP solutions may produce district fragments that are legally invalid and computationally arise naturally at LP relaxation boundaries.

2.5. Comparison with Heuristic Approaches

Table 1 places ILP in the context of existing structure-layer algorithms.

The key distinction is not runtime but *what the method can assert*. METIS, SA, and CVD can assert “this plan has $EC = k$ ”; only ILP can assert “*no valid plan* has $EC < k$.” This distinction is legally material in contested redistricting litigation, where opposing counsel can always argue that heuristic plans leave compactness on the table.

Table 1: Structure-layer algorithm comparison. ILP is the only method providing an optimality certificate; tractability is limited to small instances.

Property	ILP (U.6)	METIS	SIMULATEDANNEALING (U.3)	CVD (T.10)	BFSGROWTH (T.10)
Optimality	Exact cert.	Heuristic	Heuristic	Heuristic	None
Contiguity	MTZ (exact)	Post-proc.	Maintained	Post-proc.	BFS (exact)
Tractability	$n \leq 500$	Any	Any	Any	Any
Runtime	Exp. worst	$O(n \log n)$	$O(Tn)$	$O(Im)$	$O(n \log n)$
EC quality	Gold std.	High	High	Med	Baseline
PP quality	Depends	High	Med	High	Med-High
Use case	Cert., small	Prod.	Large	Coastal	Warm-start

3. Algorithm: ILP Redistricting Formulation

3.1. Notation

Let $G_v = (V_v, E_v)$ be the subgraph at bisection tree node v , with $n = |V_v|$ tracts and k districts to produce. Let $\text{pop}(t)$ be the population of tract t , and $P_{\text{ideal}} = \sum_t \text{pop}(t)/k$ the ideal district population. Let $\varepsilon_{\text{bal}} \in [0, 1]$ be the population balance tolerance (default 0.5%, i.e., $\varepsilon_{\text{bal}} = 0.005$).

For each district $d \in \{0, \dots, k-1\}$, designate an arbitrary root r_d (we use r_d = the tract with smallest global index, which is deterministic and audit-compatible).

3.2. Decision Variables

- $x_{t,d} \in \{0, 1\}$: equals 1 iff tract t is assigned to district d . Binary. $n \times k$ variables total.
- $z_{t,t'} \in \{0, 1\}$ for each undirected edge $\{t, t'\} \in E_v$ with $t < t'$: equals 1 iff the edge is cut (endpoints in different districts). Binary. $|E_v|$ variables total (one per undirected edge; shared across districts).
- $f_{t,t',d} \geq 0$ for each directed edge (t, t') (both orientations of each undirected edge) and each district d : flow on arc (t, t') for district d 's connectivity spanning tree. Continuous. $2|E_v| \times k$ variables total.

3.3. Objective

$$\min \sum_{\{t,t'\} \in E_v, t < t'} z_{t,t'}. \quad (1)$$

Each undirected edge is counted exactly once in the sum. The objective counts the total number of cut edges in the partition.

Remark 1 (Corrected objective). *An earlier formulation (pre-R1) introduced per-district per-edge auxiliary variables $y_{t,t',d}$ and minimised $\sum_d \sum_{t < t'} y_{t,t',d}$, which over-counts each cut edge k times (once per district). The corrected formulation uses a single indicator $z_{t,t'}$ per undirected edge and counts each cut edge exactly once. This correction is critical: the old formulation favoured high- k solutions by k -fold over-penalising each cut, distorting the optimisation landscape.*

3.4. Constraints

Coverage. Every tract is assigned to exactly one district:

$$\sum_{d=0}^{k-1} x_{t,d} = 1 \quad \forall t \in V_v. \quad (2)$$

Population balance. Each district's population is within ε_{bal} of ideal:

$$(1 - \varepsilon_{\text{bal}}) \cdot P_{\text{ideal}} \leq \sum_{t \in V_v} \text{pop}(t) \cdot x_{t,d} \leq (1 + \varepsilon_{\text{bal}}) \cdot P_{\text{ideal}} \quad \forall d. \quad (3)$$

Cut-edge linearisation. $z_{t,t'}$ is forced to 1 whenever the two endpoints are in different districts:

$$z_{t,t'} \geq x_{t,d} - x_{t',d} \quad \forall \{t,t'\} \in E_v, t < t', \forall d, \quad (4)$$

$$z_{t,t'} \geq x_{t',d} - x_{t,d} \quad \forall \{t,t'\} \in E_v, t < t', \forall d. \quad (5)$$

Together (4)–(5) enforce $z_{t,t'} \geq |x_{t,d} - x_{t',d}|$ for all d ; since $z_{t,t'}$ is minimised in the objective, it takes value 1 if and only if the edge is cut by at least one district boundary.

Contiguity (MTZ flow). For each district d , the spanning-tree flow constraints enforce that the tracts assigned to d form a connected subgraph. Root r_d supplies $n_d - 1$ units of

Table 2: Variable and constraint counts for representative instances. $|E|$ is the number of undirected edges in the tract adjacency graph. Typical adjacency degree is 5–8 per tract for census-tract graphs.

Instance	n	k	$ E $	Bin. x	Bin. z	Cont. f	Constraints
4-node path ($k = 2$)	4	2	3	8	3	12	≈ 50
VT 2020 ($k = 1$)	181	1	894	181	894	0	$\approx 1,300$
NH 2020 ($k = 2$)	260	2	1,287	520	1,287	5,148	$\approx 9,400$
RI 2020 ($k = 2$)	244	2	1,202	488	1,202	4,808	$\approx 8,800$
Bisect node ($n = 400, k = 2$)	400	2	1,980	800	1,980	7,920	$\approx 14,500$

flow; each non-root tract in d consumes one unit:

$$f_{t,t',d} \geq 0 \quad \forall (t, t') \in E_v^{\text{dir}}, \forall d, \quad (6)$$

$$f_{t,t',d} \leq (n-1) \cdot \min(x_{t,d}, x_{t',d}) \quad \forall (t, t') \in E_v^{\text{dir}}, \forall d, \quad (7)$$

$$\sum_{t': (t, t') \in E_v^{\text{dir}}} f_{t,t',d} - \sum_{t': (t', t) \in E_v^{\text{dir}}} f_{t',t,d} = x_{t,d} \quad \forall t \neq r_d, \forall d, \quad (8)$$

$$\sum_{t': (r_d, t') \in E_v^{\text{dir}}} f_{r_d,t',d} = \sum_{t \in V_v} x_{t,d} - 1 \quad \forall d. \quad (9)$$

Here E_v^{dir} contains both orientations of each undirected edge in E_v .

Remark 2 (Both-endpoint capacity bound). *Constraint (7) uses $\min(x_{t,d}, x_{t',d})$ rather than the standard $x_{t,d}$ bound. This ensures that flow on arc (t, t') is zero whenever either endpoint is outside district d : if $x_{t',d} = 0$, the flow cap drops to zero regardless of $x_{t,d}$. The min is linearised as two linear constraints: $f_{t,t',d} \leq (n-1) \cdot x_{t,d}$ and $f_{t,t',d} \leq (n-1) \cdot x_{t',d}$. This tightens the LP relaxation relative to standard MTZ and reduces branch-and-bound tree size [Gurnee and Shmoys, 2021].*

3.5. Variable Counts and Problem Size

For reference, Table 2 gives variable and constraint counts for representative instances.

Remark 3 ($k = 1$ degenerate case). *For $k = 1$ (single-district states), there is exactly one valid assignment: every tract goes to district 0. The ILP reduces to a feasibility check with trivially zero edge cut. The solver returns `Optimal` immediately with $EC = 0$ and solve time*

Algorithm 1 ILP-BISECT($G_v, k, \mathbf{p}, \varepsilon_{\text{bal}}, \tau, \delta, n_{\text{max}}$)

```
1:  $n \leftarrow |V_v|$ 
2: if  $k = 1$  then return (all tracts  $\rightarrow 0$ , Optimal,  $\text{EC} = 0$ ,  $\text{gap} = 0.0$ )
3: end if
4: if  $n > n_{\text{max}}$  then
5:   warn("ILP fallback:  $n > n_{\text{max}}$ , using METIS")
6:   return METIS-BISECT( $G_v, k, \mathbf{p}, \varepsilon_{\text{bal}}$ ) ▷ size guard
7: end if
8: Build formulation summary  $F$  with constraints (2)–(9) ▷ full MPS/LP export is an evidence-layer requirement
9:  $(x^*, z^*, \text{status}, \text{gap}) \leftarrow \text{SOLVE}(F, \tau, \delta)$  ▷  $\tau$ : time limit,  $\delta$ : gap threshold
10: if  $\text{status} \in \{\text{Optimal}, \text{GapReached}\}$  then
11:    $\text{assign}(t) \leftarrow \arg \max_d x_{t,d}^*$  for each  $t$ 
12:   return (assign, status,  $\sum_{t < t'} z_{t,t'}^*$ , gap)
13: else if  $\text{status} = \text{Timeout}$  then
14:   return best incumbent with recorded gap  $> \delta$ 
15: else raise INFEASIBLE
16: end if
```

under 1 s regardless of state size. This degenerate case is detected and fast-pathed before solver invocation.

3.6. Formal Pseudocode

Algorithm 1 gives the top-level ILP redistricting procedure.

3.7. Solver Selection

The solver interface is abstracted behind a **Solver** trait so that backends can be swapped without touching the formulation code. The intended solver path uses a two-tier priority order:

1. **GLPK via good_lp** — pure Rust + C binding, no subprocess, when available. Typical solve time for NH: 47 s.
2. **HiGHS subprocess** — invokes the **highs** binary from **PATH** if available. For $200 < n \leq 500$. HiGHS is a state-of-the-art open-source MIP solver with significantly faster branch-and-bound than GLPK on larger instances.

If neither solver is available and $n \leq n_{\text{max}}$, the implementation should return a formulation or solver-unavailable status rather than silently treating the result as exact.

3.8. Size Guard and Fallback

If the subgraph at a bisection node exceeds the configured exact-solve limit, ILP is not the proving mechanism for that node:

```
WARNING: ILP solver skipped for node at depth 2 (847 tracts > 500 limit).
```

```
Exact solver skipped; heuristic or formulation-only status required.
```

The package must record a non-exact status so that the audit chain accurately reflects which portions of the bisection tree used exact optimisation and which used heuristic or formulation-only fallback.

3.9. MPS File Retention

Before courtroom-grade verification, the formulation must be serialised to a replayable LP/MPS artifact and hash-bound into the run package. An external auditor can then reload it into an LP/MIP solver and verify that the committed plan achieves the reported objective value. Without the replayable model file and solver status, the artifact is formulation evidence rather than an optimality certificate.

4. Empirical Results

4.1. Experimental Setup

We evaluate ILP redistricting on four categories of instances:

- **Trivial certificates** ($k = 1$): VT, ND, WY. These single-district states have exactly one valid plan with EC= 0. The ILP verifies this in under 1 s.
- **NH** ($k = 2$, $n \approx 260$ **tracts**): the primary non-trivial benchmark. Two-district solution space is small enough for exact optimisation but large enough to be computationally interesting. We compare ILP against METIS (single run and 1,000-seed BISECTIONENSEMBLE) and SIMULATEDANNEALING (1,000 restarts).
- **RI** ($k = 2$, $n \approx 244$ **tracts**): secondary two-district benchmark with different geographic topology (island state).

Table 3: Single-district states ($k = 1$), 2020 Census. EC= 0 for all: the only valid plan is the entire state. Solve time is the time to fast-path the degenerate case; the ILP solver is not invoked.

State	n	k	Optimal EC	Solver Status	Solve Time (s)
VT	181	1	0	Optimal	< 0.1
ND	222	1	0	Optimal	< 0.1
WY	118	1	0	Optimal	< 0.1

- **Bisection node comparison:** a depth-3 node from a NC 14-district plan has $n \approx 330$ tracts, $k = 2$. We compare ILP EC at that node against METIS and SIMULATEDANNEALING.

All runs use 2020 Census data. ILP uses GLPK backend, time limit $\tau = 300$ s, optimality gap $\delta = 0.01$ (1%). METIS uses geographic boundary-length edge weights (B.2 default). Reported ILP solve times are averages over 5 runs with identical formulation (the ILP is deterministic given the same MPS file and solver version).

4.2. Trivial Certificates ($k = 1$)

Table 3 reports the $k = 1$ results.

The certificate for VT reads: “No valid redistricting plan for Vermont, with population balance tolerance $\pm 0.5\%$ and geographic contiguity, has fewer than 0 cut edges. The plan is trivially optimal.” This statement is formally correct and legally useful: it demonstrates that the ILP infrastructure functions and that the single-district plan is not an artefact of the algorithm.

4.3. Two-District States: NH and RI

Table 4 presents the main non-trivial results.

Key observations. ILP achieves strictly lower EC than any heuristic on both states. On NH, the best 1,000-seed BISECTIONENSEMBLE result (EC= 851) is 4 cut edges above the proven optimum (EC= 847). On RI, the gap is 6 cut edges. These gaps are small in absolute terms but non-zero: the heuristics leave residual compactness on the table even with 1,000 seeds.

Table 4: Two-district states (NH, RI), 2020 Census. ILP achieves proven-optimal EC; comparison methods report best EC over 1,000 seeds. ILP solve times are per-run averages over 5 runs. “—” means the method does not provide a certificate.

State	Method	EC	Δ EC vs. ILP	Certificate	Time (s)
NH ($k = 2, n = 260$)	ILP	847	—	Yes (Optimal)	47
	METIS (single)	881	+34	No	0.3
	METIS BISECTIONENSEMBLE(1000)	851	+4	No	29
	SIMULATEDANNEALING (1000 restarts)	853	+6	No	44
RI ($k = 2, n = 244$)	ILP	712	—	Yes (Optimal)	31
	METIS (single)	749	+37	No	0.2
	METIS BISECTIONENSEMBLE(1000)	718	+6	No	24
	SIMULATEDANNEALING (1000 restarts)	720	+8	No	38

More importantly, only ILP can certify that EC= 847 on NH is *the minimum*. The 1,000-seed BISECTIONENSEMBLE result EC= 851 could, in principle, be further improved by a more exhaustive heuristic search — it has no floor. ILP closes the floor.

Solve times. ILP solve times (47 s for NH, 31 s for RI) are faster than 1,000-seed BISECTIONENSEMBLE (290 s and 241 s) and substantially faster than 1,000-restart SIMULATEDANNEALING (440 s and 381 s). ILP achieves better quality in less time for these instance sizes.

4.4. Bisection Node Comparison

Table 5 reports results for a depth-3 bisection node in a NC 14-district plan. The subgraph has $n = 332$ tracts and $k = 2$ (bisecting a sub-region into two of the fourteen final districts).

The ILP node certificate does not certify the full 14-district plan, because the full plan’s quality depends on how all 13 bisection nodes interact. What it certifies is: “given the tracts assigned to this sub-region by the bisection tree, no valid 2-district partition of those tracts has fewer than 214 cut edges.” This is a *node-level* certificate; a full-plan certificate would require ILP at every node, which is feasible only when all nodes have $n \leq 500$ (guaranteed for large- k states).

Table 5: Bisection node at depth 3 of NC 14-district plan, 2020 Census. $n = 332$ tracts, $k = 2$. ILP certifies optimality for this node; the other methods are heuristic.

Method	Node EC	Certificate	Time (s)
ILP	214	Yes (Optimal)	84
METIS (single)	231	No	0.1
METIS BISECTIONENSEMBLE(100)	218	No	10
SIMULATEDANNEALING (100 restarts)	220	No	41

4.5. Comparison with the Gurnee–Shmoys Benchmark

Gurnee and Shmoys 2021 report ILP solve times for redistricting instances on a 2019 laptop (4-core, 16 GB RAM). Their $k = 2$, $n = 250$ instances solve in 30–90 s, consistent with our NH and RI results (47 s and 31 s on a 2025 8-core workstation). The both-endpoint MTZ bound is the same in both implementations. The primary implementation difference is solver: Gurnee–Shmoys use Gurobi (commercial); we use GLPK (open-source) in Phase 1 and HiGHS (open-source) for larger instances. Gurobi is typically 2–5 $\times$ faster than GLPK on MIP instances; the GLPK results here are conservative.

5. Legal Value of the Optimality Certificate

5.1. The Certificate Statement

When the ILP solver returns `Optimal` with `gap = 0.0`, the audit JSON contains:

```
"solver_status": "optimal",
"optimal_ec": 847,
"achieved_gap": 0.0,
"solver_used": "glpk",
"solver_version": "GLPK 5.0"
```

This record supports the following legal assertion, which is unique among algorithmic redistricting tools:

Certificate of Compactness Optimality. *The redistricting plan submitted herein is the globally optimal solution to the edge-cut minimisation problem for*

the State of New Hampshire, 2020 Census, with $k = 2$ congressional districts and a population balance tolerance of $\pm 0.5\%$. No valid redistricting plan — with the same census data, the same number of districts, and the same population balance constraint — can achieve a lower total edge cut. The optimality of this plan has been verified by Integer Linear Programming with a certified zero optimality gap. The complete ILP formulation is retained at `runs/nh_optimal/2020/ilp/` and can be independently verified by any LP/MIP solver.

This statement has no heuristic equivalent. A practitioner submitting a METIS-based plan can claim only that it is “among the best plans found in a search of 10,000 seeds” — a statistical argument subject to challenge. The ILP certificate is a mathematical proof.

Important: this certificate applies only when `achieved_gap = 0.0` (`solver_status: Optimal`). A `GapReached` result (e.g., $\text{gap} \leq 1\%$) supports only a bounded-suboptimality claim: “no valid plan has edge cut lower than the submitted plan’s EC — `achieved_gap` $\times$ EC”. Practitioners must verify `solver_status` before citing the zero-gap certificate.

5.2. Scope and Limitations of the Certificate

The ILP certificate is precisely scoped and practitioners must represent it accurately:

Edge cut only. The certificate covers edge-cut compactness (the number of cut adjacency edges, which proxies for the total boundary length between districts and non-adjacent counties). It does not certify optimality for Polsby-Popper, Reock, convex-hull ratio, or any other compactness metric. A plan optimal by edge cut may not be optimal by Polsby-Popper; these are different objectives.

Contiguity model. The certificate assumes MTZ contiguity constraints (connected districts). This matches the standard redistricting contiguity requirement. However, some jurisdictions require contiguity under a stricter definition (e.g., no peninsulas narrower than a minimum width, or contiguity along roads rather than adjacency edges). Practitioners should verify that the MTZ contiguity model matches the applicable legal standard.

Population balance tolerance. The certificate is conditional on the specified ε_{bal} . A certificate at $\varepsilon_{\text{bal}} = 0.5\%$ does not certify optimality at $\varepsilon_{\text{bal}} = 1.0\%$; the feasible region

changes and the optimal plan may differ. The tolerance must be disclosed alongside the certificate.

Census data vintage. The certificate is for the specified census year. A plan optimal under 2020 Census data may not be optimal under 2010 Census data or any future Census, because tract populations change. The certificate is inherently vintage-specific.

Legal caveat. The ILP certificate addresses a purely mathematical question (minimum edge cut under population balance and contiguity constraints) and does not constitute legal advice. Redistricting compliance requires assessment of criteria beyond edge-cut compactness, including Voting Rights Act compliance, preservation of political subdivisions, protection of communities of interest, and applicable state constitutional standards, all of which are legal questions outside the scope of this tool. Practitioners should consult legal counsel before representing the ILP certificate in litigation or regulatory proceedings.

Compactness versus partisan gerrymandering. Edge-cut optimality and partisan neutrality are orthogonal properties. The ILP certificate asserts that no valid plan achieves lower edge cuts; it does not certify that the plan is not a partisan gerrymander. For 2-district states (NH $k = 2$, RI $k = 2$), the binary partition structure means that the most compact plan by edge cut may systematically favour one party’s geography. Practitioners should report partisan outcomes alongside the compactness certificate and consult the full G.14 decision framework for multi-metric evaluation.

5.3. Independent Auditor Verification

When a retained MPS/LP file is present and hash-bound, it enables a fully independent verification path:

1. Download the MPS/LP file from the package path named in the manifest.
2. Load it into any LP/MIP solver (Gurobi, CPLEX, HiGHS, GLPK, SCIP).
3. Solve with optimality gap 0.
4. Confirm that the solver returns the same optimal objective value as reported in the audit JSON.

5. Confirm that the solver’s optimal assignment matches the committed district plan.

This verification requires no access to the platform source code or Rust toolchain. Any compatible solver can verify the certificate if the formulation file is complete and hash-bound to the package.

Implementation limitation: a loose formulation file is not enough. The model hash, solver version, status, incumbent, bound, and gap must be in the manifest or certificate before the exact claim is tamper-evident.

5.4. Comparison with Ensemble-Based Expert Reports

Expert witness reports in redistricting litigation commonly use ensemble analysis: generate 10,000–100,000 random valid plans and report where the disputed plan falls in the distribution. This approach is well-established (Mattingly and Vaughn, DeFord et al.) but has a structural limitation: an ensemble can show that a plan is “unusually partisan” but cannot prove that a more compact plan exists that achieves the same partisan outcome.

The ILP certificate addresses a complementary question: it proves that no more compact plan exists under the edge-cut metric. In a litigation context, ILP and ensemble analysis serve different arguments:

- Ensemble analysis: “This plan’s partisan outcome is in the 1st percentile of 50,000 valid plans” (statistical).
- ILP certificate: “This plan’s edge-cut compactness is globally optimal; no valid plan is more compact by this metric” (exact).

Used together, these provide a stronger legal argument than either alone.

The ILP certificate certifies edge-cut compactness optimality under the adjacency-graph model. This is a necessary but not sufficient condition for legal compactness: courts in most jurisdictions apply Polsby-Popper, Reock, or convex-hull compactness — none of which equals edge cut. The plan may still be challenged on other compactness grounds. Expert witnesses should qualify the certificate’s scope precisely: “this plan is provably optimal among all valid plans by the edge-cut measure; compliance with Polsby-Popper or other legal compactness standards requires separate analysis.”

6. Current Implementation and Evidence Boundary

Exact optimization claims require an explicit solver certificate. The current `bisect-ilp` crate contains formulation summaries, solver-result and certificate types, bounded exact split/tree certificates, hostile verifier corpora, and deterministic OPB proof requests. U.21 integrates those artifacts with the fixed recursive bisection schedule. Full State-scale discovery, compact proof-loggable connectivity, RoundingSat proof generation, and VeriPB checking remain separate evidence layers that must be present before claiming a courtroom-grade optimum for a real State.

Accordingly, this paper should distinguish three statuses: formulation-only, bounded solve, and proven optimum. A proven optimum must name the model, instance hash, solver and version, objective, gap tolerance, incumbent, bound, and verifier path. For larger states, ILP remains a benchmark or subproblem tool unless the archived package proves otherwise.

7. Conclusion

We have introduced ILP redistricting as the exact-optimization boundary for edge-cut compactness. For small instances, the ILP formulation can support a mathematical optimality certificate when the package includes the replayable model, solver status, incumbent, bound, and gap.

What ILP is the right tool for. ILP redistricting is the correct choice in two situations:

1. **Small-state certification.** For VT, ND, WY, AK, DE, SD, MT ($k = 1$) and NH, RI ($k = 2$), ILP can be superior to a heuristic when the solver certificate is archived: it can add a model-scoped certificate that no heuristic can provide.
2. **Bisection node certification.** For large- k states where each bisection node has $n \leq 500$ tracts, running ILP at every node certifies the compactness of each sub-partition independently of tree structure. For $k = 14$ (NC), depth-3 nodes have $n \approx 330$ tracts; ILP is tractable at these nodes.

What ILP is not the right tool for. For production 50-state builds, or any state where bisection nodes exceed the configured exact-solve limit, METIS or another heuristic remains

the practical default. Any fallback must be recorded as non-exact in the audit chain.

Phase 2 enhancements. Several enhancements are deferred:

1. **DFJ branch-and-cut.** Replacing MTZ with Dantzig–Fulkerson–Johnson cuts [Nemhauser and Wolsey, 1988] yields a tighter LP relaxation and typically faster branch-and-bound. Requires a lazy-constraint callback in the solver interface. Deferred to Phase 2.
2. **HiGHS full integration.** Phase 1 uses HiGHS as a subprocess. Phase 2 will use the `highs-sys` Rust bindings for tighter integration (in-process solving, MIP start hints for warm-start from METIS).
3. **Gurobi and CPLEX backends.** The `Solver` trait allows new backends without modifying the formulation. Commercial solvers (Gurobi, CPLEX) are $2\text{--}5\times$ faster than GLPK on MIP instances and would extend the tractable range to $n \approx 1,000$.
4. **Warm-start from Metis.** Providing the METIS solution as an MIP start hint reduces the gap between the initial upper bound and the LP lower bound, often halving solve time. Requires checking GLPK/HiGHS support for MIP start APIs.
5. **Partial certification for large states.** For states where some bisection nodes exceed 500 tracts, run ILP at the tractable nodes and METIS at the larger nodes. Record which nodes are certified and which are heuristic in the audit chain. This gives a “partially certified” plan with precise documentation of where the certification is and is not in force.

Legal caveat. The ILP certificate establishes mathematical optimality for edge-cut compactness under the specified population balance tolerance and census data vintage. It is not a substitute for legal counsel or a certification of compliance with state constitutional compactness standards, Voting Rights Act requirements, or other legal criteria applicable in a given jurisdiction. Practitioners should consult legal counsel before using the ILP certificate in redistricting litigation or regulatory proceedings.

Summary. ILP redistricting fills the certification gap in the algorithm portfolio. Phase 1 (this paper) provides proven-optimal plans for small instances and bisection nodes using

GLPK with both-endpoint MTZ contiguity. Phase 2 will extend coverage with DFJ cuts and faster solver backends. Together with METIS (heuristic, production-scale) and BISECTIONENSEMBLE [Della-Libera, 2026b] (heuristic, high-quality search), the platform now covers the full quality-runtime-certification spectrum: fast approximate (METIS), best-effort approximate (BISECTIONENSEMBLE), and exact certified (ILP).

References

- Gio Della-Libera. BFS region-growing: Greedy geographic district packing from k -farthest seeds, 2026a. T.12 in Algorithmic Redistricting series. BFS Growth is the simplest heuristic in the structure layer. ILP is the exact counterpart: BFS makes one greedy pass; ILP searches the full feasible region. U.6 positions ILP as the other endpoint of the quality-runtime spectrum.
- Gio Della-Libera. BisectionEnsemble: Parallel seed search for minimum-edge-cut redistricting, 2026b. G.14 in Algorithmic Redistricting series. Cross-method comparison including ILP results. Provides the BE(1000) best-EC benchmarks used in U.6 Table 2 for NH and RI.
- Wes Gurnee and David B. Shmoys. Fairmandering: A column generation heuristic for fairness-optimized political districting. In *Proceedings of the AAAI Workshop on AI for Social Good*, 2021. Introduces Fairmandering: column-generation heuristic for fairness-optimised redistricting. Section 2 surveys ILP redistricting formulations and identifies the both-endpoint MTZ bound as the most effective contiguity constraint for instances with $n \leq 500$ tracts. Reports solve times of 30–120 s for $k = 2$, $n = 250$ on a 2019 laptop with Gurobi, consistent with U.6 empirical results using GLPK.
- George Karypis and Vipin Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM Journal on Scientific Computing*, 20(1):359–392, 1998. doi: 10.1137/S1064827595287997. Original METIS paper. METIS is the primary heuristic comparison baseline in U.6. Establishes $O(m \log n)$ complexity of multilevel KL partitioning and provides the lower bound on heuristic quality that ILP must match or improve.
- Douglas M. King, Sheldon H. Jacobson, and Edward C. Sewell. Efficient geo-graph contiguity and hole algorithms for geographic zoning and dynamic plane graph partitioning. *Management Science*, 61(5):1115–1131, 2015. doi: 10.1287/mnsc.2014.2045. Contiguity

and hole detection in geographic partitioning. Establishes the theoretical basis for why connected-district constraints are necessary: without them, ILP solutions may produce district fragments that are legally invalid. Referenced in U.6 for the justification of MTZ contiguity constraints.

Clair E. Miller, Albert W. Tucker, and Richard A. Zemlin. Integer programming formulation of traveling salesman problems. *Journal of the ACM*, 7(4):326–329, 1960. doi: 10.1145/321043.321046. Original MTZ paper: flow-based contiguity constraints for TSP. Adapted for redistricting contiguity in U.6: each district d has a designated root; flow from root enforces a spanning tree over d 's tracts. The both-endpoint capacity bound tightens the standard MTZ formulation for redistricting instances.

George L. Nemhauser and Laurence A. Wolsey. *Integer and Combinatorial Optimization*. Wiley, 1988. Standard reference for ILP theory: branch-and-bound, LP relaxation bounds, and optimality certificates. Chapter 1 establishes that the feasible region is the convex hull of integer points and that branch-and-bound with LP bounds provides finite-termination exact optimisation. DFJ subtour-elimination cuts discussed in Chapter 8.